

Simultaneous Personnel Localization and Mapping



WPI

A Major Qualifying Project Report Submitted to the Faculty of
Worcester Polytechnic Institute

In partial fulfillment of the requirements for the
Degree of Bachelor of Science in Electrical & Computer Engineering

By:

Georges Gauthier
John DeCusati

October 29, 2016

Advisor:
Professor R. James Duckworth

Contents

Abstract	i
1 Introduction	1
2 Background	3
3 Methodology	9
3.1 Camera Testing	9
3.1.1 Camera Operation	9
3.1.2 I ² C Control	12
3.1.3 Data Management	15
3.1.4 Transmitting Images Over UART for Analysis	19
3.2 Final Camera Hardware Implementation	22
3.3 Disparity Algorithm	25
3.3.1 Sum of Absolute Differences	26
3.3.2 Image Rectification	29
3.3.3 Test Implementation	31
3.3.4 Final Implementation	33
References	34
4 Appendix	35
4.1 Useful Resources	35
4.2 Component Selection	37
4.3 Camera Module Decision Matrix	37
4.4 Camera Module Control Register	39
4.5 Stereo Camera Schematic	40
4.6 MATLAB Code	41

4.6.1	Camera Image Parsing	41
4.6.2	Disparity Algorithm Implementation	42
4.7	Verilog Code	46
4.7.1	MT9V034 and Al422b Test Code	46
4.8	C Code	53
4.8.1	Camera Image Parsing	53
4.9	LaTeX Coding Examples	56
4.9.1	Figures	56
4.9.2	Code Snippet	56
4.9.3	Using the bibliography	57

List of Figures

1	Real-Time SLAM with a Single Camera [3]	3
2	Serveball's Squito [6]	4
3	Serveball's Squito Input and Output [6]	5
4	From Left to Right: Original Image, Disparity Map, Object Detection Results [7] . .	5
5	Human Body Tracking by Adaptive Background Models and Mean-Shift Analysis [12]	6
6	Functional Block Diagram	7
7	LI-VM34LP Breakout Board	10
8	Frame and Line Valid [11]	10
9	Line Data Transfer [11]	11
10	Camera Data Transfer	12
11	Example I ² C Data Transfer	13
12	Example I ² C Transfer with Camera	14
13	Camera Trigger and FV in Trigger Mode	15
14	Camera Test System Block Diagram	18
15	Transferring Line Data from FIFO to FPGA	19
16	Reading FIFO Data	20
17	Notebook With Grid and Oscilloscope Leads	20
18	Camera Test Setup	21
19	Stereo Camera PMOD PCB	24
20	Stereo Camera Breakout Under Test	24
21	Stereo Camera Breakout Sample Image	25
22	Sum of Absolute Differences [9]	27
23	Block Matching Overview [2]	28
24	Disparity Algorithm Output	29
25	Horizontal Epipolar Lines [2]	30
26	Stereo Image Rectification [8]	31
27	Disparity Test Implementation	32

28	Disparity Final Implementation	33
29	A Test Figure	56

Abstract

The overall goal of this project is to create a device capable of generating detailed maps and imagery of an area in real-time. This device will rely on an FPGA, and will use image processing algorithms capable of detecting, localizing, and tracking human beings. The device will gather imagery from the visual light and infrared-spectrums, as well as localization data and distance measurements from an IMU and a rangefinder. Along with gathering and processing data, the device will also serve as a long-range wireless access point, and will be able to transmit all generated maps and imagery in real-time. A major deliverable of this project is that the transmitted data will be fully processed, allowing it to be viewed remotely on less-powerful, mobile devices.

This device will be especially useful for first responders. It is intended to be mounted on a small remote control vehicle, allowing any connected user to wirelessly traverse dangerous and remote locations in search of people in need. Since this device transmits data in real-time, it will be able to provide first responders with an accurate representation of not only a 2-D floor plan of an area such as a building, but also where any people are located. An anticipated use of this device would be in the event of a building in danger of collapsing. Since it would be dangerous to physically enter the building, first responders could locate any people trapped inside and find the fastest route to them using the wirelessly transmitted floorplan. The first responders would also be aware of any dangers in their way by making use of the real-time augmented video stream. This video stream will consist of image data with overlaid with object indicators and location information on any human beings detected by the image processing algorithms.

1 Introduction

Currently there are many applications that rely on a simple video camera setup in order to gather information on remote and inaccessible locations. Although this is an effective strategy for simple surveillance, it is limited in many ways. Using current imaging and sensor technology, it is possible to gather camera images and 3D depth information on a given area as both a cost-effective and information-rich alternative to using a camera module on a device. If a product were to be created that gathers this information as a replacement to using a standalone camera module, it would be possible to use high speed data processing techniques in both hardware and firmware that would allow for the creation of an augmented real-time video feed.

This type of technology is known as Simultaneous Localization And Mapping, or SLAM. The purpose of SLAM is to compute the location of an agent within its environment, and allows for the creation of self-aware robot systems that are able to respond to their surroundings. SLAM is a common area of research in the image processing and high-speed computing field, and has been applied mainly to autonomous vehicles. We would like to propose the creation of a SLAM-like system that is capable of monitoring and mapping its environment in real-time, as well as detecting and localizing objects, such as human beings. For the purposes of our project, we would like to define our desired objective as Simultaneous Personnel Localization And Mapping, or SPLAM.

A device that is capable of both mapping its surroundings using SLAM and performing human detection in real-time would be applicable to many different fields. We are especially interested in creating a proof of concept sensor suite capable of performing these tasks that can be added to existing robotic systems as a stand-in replacement for a video camera. This type of technology would allow for people such as firefighters or first responders to wirelessly traverse dangerous and remote locations in search of people in need. We envision our sensor suite being able to process data so that its users would be provided with a 2D “floorplan” of the area being traversed by the sensor suite, as well as an augmented video feed with

imagery containing indicators for any detected human beings in the area.

One type of technology that would be useful for performing the high speed data processing necessary for SPLAM is a Field Programmable Gate Array , or FPGA. FPGAs pose several advantages over using standard computing or microcontroller technology for real-time data processing, as they have the ability to manipulate digital information in parallel using hardware only. This allows for extremely high-speed performance, as calculations can be run in parallel and are only dependent on their data inputs as opposed to waiting for specific tasks or scripts to run on a microcontroller or computer software interface.

Although FPGA technology is highly applicable to performing SLAM-like tasks due to its high speed, there are currently few existing commercial products that use FPGAs for the purpose of performing SLAM. Most current SLAM implementations rely on the use of a sensor suite connected to a computer or system on chip (SoC) computing device that performs data analysis using software or a real-time operating system (RTOS). This means that data must first be collected by a sensor suite, and then transferred to an external computing device that is only capable of processing it serially based on its arrival time. Although this type of setup is acceptable for performing real-time situational awareness analysis, we propose that an embedded, FPGA-based SPLAM device would be a much more elegant and higher-speed solution.

2 Background

A major concern with real-time image processing, especially in first responder situations, is speed. Because FPGAs have the ability to process data in parallel, they are ideal for this type of application. Using an FPGA for this system will enable all data inputs to be processed at the same time, thereby dramatically increasing throughput speed. Dealing with each input separately makes it easier to combine everything together, especially because each component functions at a different clock speed. Also, since everything is running in parallel, more cameras can be added to the system to increase the field of vision of the device without introducing any latency in the system, as long as enough memory is available.

SLAM is a widely expanding field with much potential for improvement. One application of such a system is a proof of concept of camera-based SLAM systems, presented by Andrew Davison of Oxford University in a research paper entitled "Real-Time SLAM with a Single Camera" [3]. This system is handheld and relies on a computer using a 2.2 GHz Pentium processor connected to a single camera and laser rangefinder. The system requires prior knowledge of the area being analyzed before it can successfully localize and map. It implements edge detection, but is limited to the narrow field of vision of the rangefinder, so it is only able to map an object directly in front of it. This system carries a latency of around 33 milliseconds. An output frame of the device is shown in Figure 1.

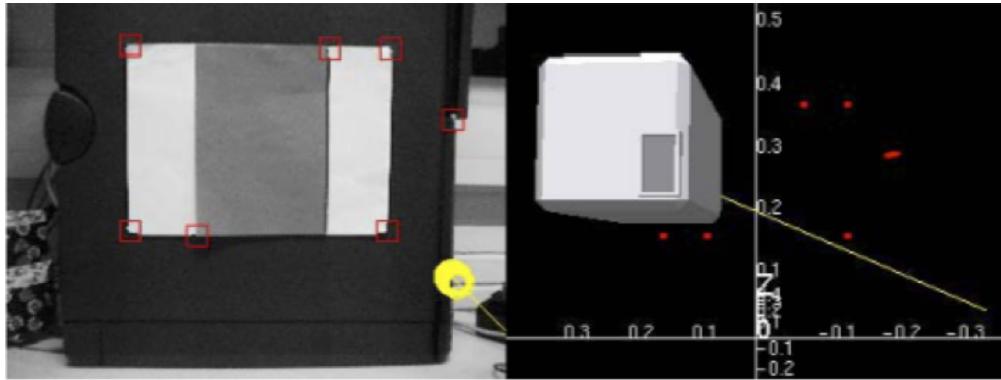


Figure 1: Real-Time SLAM with a Single Camera [3]

The frame on the left in Figure 1 is the video feed with 6 points of a paper target input as

prior knowledge, along with successfully marked identifying features (marked as red squares), and another identifying feature that is not marked for measurement (marked by a yellow circle). The frame on the right is a localization graph displaying the positions of all red squares.

A more commercial device similar to our concept is Serveball's SquitoTM [6]. Squito is a wireless, throwable, 360° panoramic camera that implements target detection to stabilize the video feed from its many cameras. It is shown in Figure 2 below.



Figure 2: Serveball's Squito [6]

Squito utilizes a microprocessor receiving input from cameras, as well as orientation and position sensors in order to transmit a real-time stabilized video of its adventure. The device is still in the prototype stage and is receiving interest from first responders. The image in Figure 3 shows the input from the Squito's four cameras on the left, and the corresponding stitched output on the right.

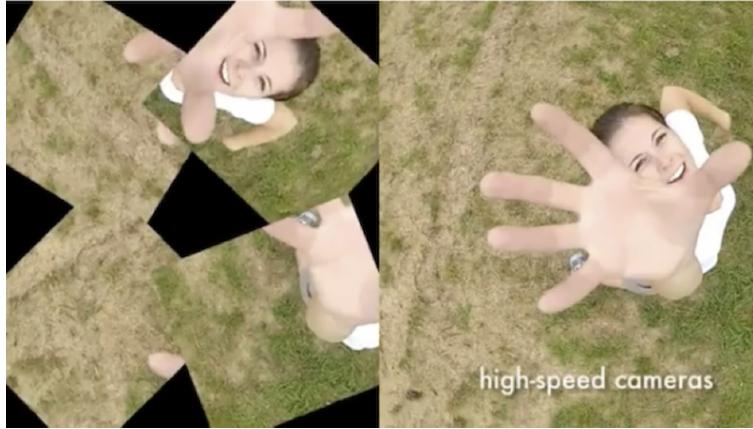


Figure 3: Serveball’s Squito Input and Output [6]

By using multiple camera sensors in a sensor suite, it is also possible to determine depth information from corresponding images of an area. This technique is known as stereo imaging, and the process of gathering depth information from a pair of stereo images is known as disparity mapping. University of Bologna researchers Stefano Mattoccia and Matteo Poggi have worked to implement a real-time disparity mapping algorithm on an FPGA, and an example of a stereo image disparity is shown in Figure 4 below [7]. Using their stereo vision algorithm, the researchers are able to generate real-time image data showing the relative locations of objects within an image frame using color gradients. Based on this depth information, it is also possible to detect objects located within the field of view of the stereo imaging system, as shown in Figure 4. An implementation similar to this would be extremely useful in a SLAM-like system, as it would allow for the localization of objects and creation of 2D “floorplans” of an area in real-time using only two camera sensors.

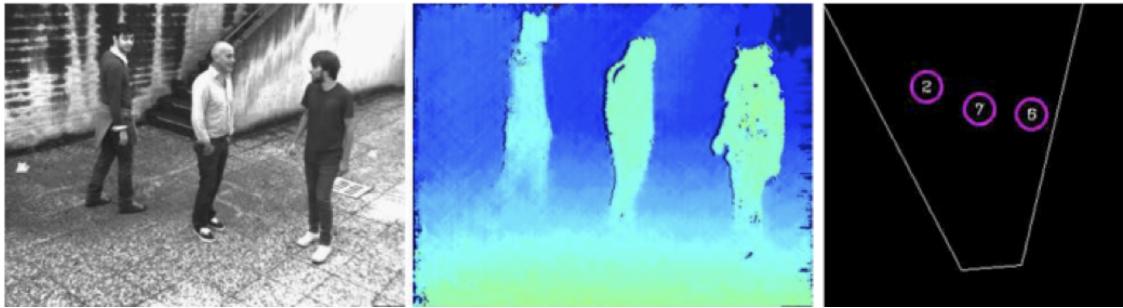


Figure 4: From Left to Right: Original Image, Disparity Map, Object Detection Results [7]

Many security systems implement human detection and human body tracking in order to increase their effectiveness. These devices process real-time images in order to identify human characteristics, and are limited to the field of vision of a stationary or rotating camera. An example of this type of system is explored by the Mitsubishi corporation in a research paper entitled “Human Body Tracking by Adaptive Background Models and Mean-Shift Analysis” [12]. The paper explores a stationary image processing system implemented on a PC platform with a 1.8GHz processor that yields a maximum processing time of 100 milliseconds. An output frame of the system is shown in Figure 5 below.



Figure 5: Human Body Tracking by Adaptive Background Models and Mean-Shift Analysis [12]

Our proposed device will combine the ideas of the four systems examined. It will be able to simultaneously localize and map an area, as well as implement human detection algorithms. The device will be capable of generating real-time 2D maps of any area it has traversed with humans' locations labeled, and an augmented video feed of what the device is recording with any humans' positions marked.

In order to successfully implement this system, we propose the creation of a device that will rely on two stereo cameras, a laser rangefinder, and an inertial measurement unit (IMU) as its sensor suite, as shown in Appendix Item 2. Limitations of previous art have been in their ability to combine human detection with real-time localization and mapping of a

large field of vision. Little to no existing commercial products are also capable of processing their gathered data locally and in real-time, with their gathered data usually requiring post-processing on external computing devices. Stereo cameras will allow our device to calculate disparity, just as human eyes do. Although disparity is useful for localization, it is not enough for accurate mapping because it only accurately provides the relative distance between objects. The inclusion of a rangefinder will allow for precise base distance readings, and an IMU will be used to spatially reference all gathered data. All of this data will be combined with the disparity maps and image data in order to create flawless localization and mapping. All time-dependent processing required for the device will be mainly done in parallel using hardware on an FPGA. An overall functional block diagram of our intended implementation is shown in Figure 6 below.

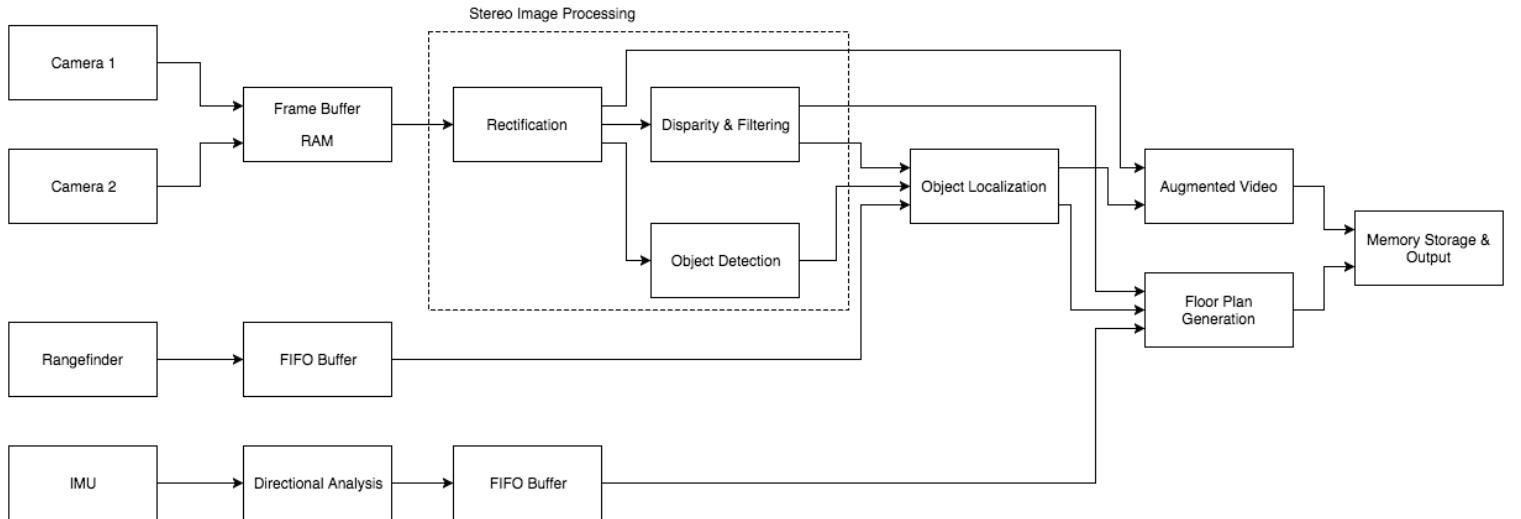


Figure 6: Functional Block Diagram

Most applicable previous camera-based systems have also focused on object detection from a stationary point, or edge detection from a mobile platform. Our project aims to combine these concepts, by creating a mobile device that detects people which will be especially of use in many first responder situations. In addition, this device will receive data from the visible light and infrared spectrums in order to identify people quickly and accurately in a way that has not been previously implemented.

As our research has progressed over time, our project objectives have continually evolved. We originally envisioned the creation of a device that used laser rangefinders to create 3D maps of its surroundings, similar to that of a Carnegie Mellon University device created in order to volumetrically map abandoned mines [13].

As our research progressed, we believed that we could use a visual light and thermal imaging camera set to gather information on an area, and supplement that data with IMU and rangefinder readings in order to produce detailed maps of our sensor suites? surroundings. Eventually we came upon the concept of disparity mapping and generating depth information from image data, and decided that we would again like to shift the overall setup of our device to rely mainly on stereo image data. Due to our overall budget and the resources that have been made available to us, in the coming terms we plan to use an electronic rangefinder, IMU, and stereo camera pair to generate real-time SPLAM video and floorplan information. Although we were also originally planning on including a thermal camera in our sensor suite as well, we have decided to eliminate the module in favor of higher quality cameras due to its prohibitive cost, low resolution, slow sampling rate, and small field of view. More information on this decision can be found in Appendix item 4.3.

3 Methodology

3.1 Camera Testing

After obtaining two of the MT9V034 cameras chosen through the process referenced in Appendix item 4.3, several steps were taken to obtain test images from each camera. These steps are outlined in the following sections.

3.1.1 Camera Operation

In order to gather working images from each camera module, we first needed to understand what circuitry our camera module breakouts contained so that we could interface with them. The MT9V034 camera breakouts used have been purchased through Leopard Imaging Inc. Although these camera module breakout boards are intended to be used with Leopard Imaging's LeopardBoard ARM development board, the breakouts were found to contain only the supporting circuitry recommended in the MT9V034 datasheet, and we decided that they would be ideal for our application [5, 11].

Once the schematics of each camera module breakout were known, it was then possible to design a basic control interface for each camera. According to the MT9V034 datasheet, each camera module needs to be supplied with an external Master Clock and Output Enable signal in order to operate [11]. A simple Verilog module for the Nexys3 Spartan-6 FPGA board was created in order to supply the camera module with a 24MHz master clock signal, and a switch was used to toggle output enable. With this module implemented, the camera module's default outputs could then be observed. In order to interface the camera module with an FPGA, the breakout board shown in Figure 7 was also created to make the module's pins more easily accessible.

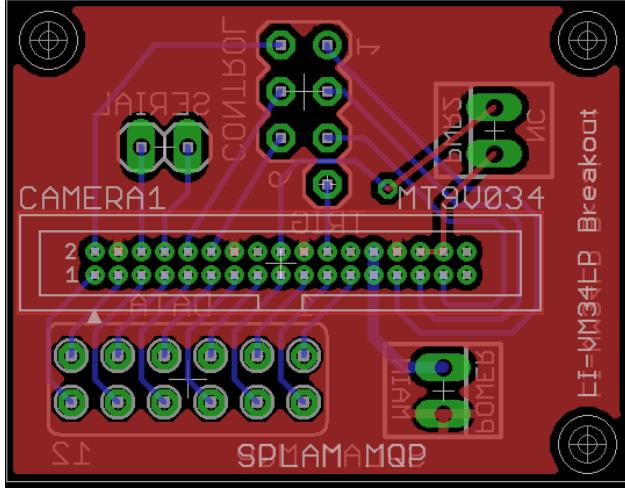


Figure 7: LI-VM34LP Breakout Board

By default, the MT9V034 camera module will continuously gather image data at 60Hz as long as it is supplied with an external clock signal and output is enabled [11]. Several output signals from the camera module are then used to transmit image data. Each image, or frame, is broken up into individual “lines” which correspond to a line of pixels that stretch the width of the frame. Since our camera module captures images at 752x480 pixel resolution, one frame will contain 480 lines of 752 pixels each. The camera module breaks up image data by frame and line, and camera data pins FRAME_VALID and LINE_VALID are toggled to indicate the transmission of a frame or line. The timing diagram shown in Figure 8 shows the operation of these pins while transmitting an image.

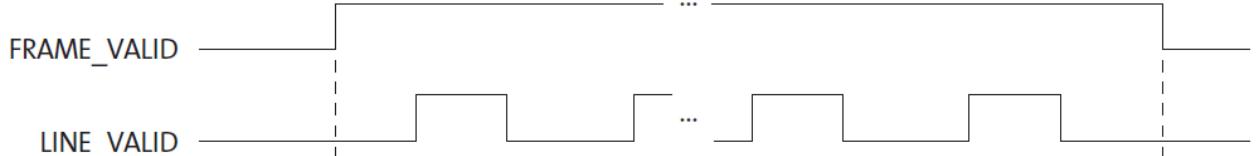


Figure 8: Frame and Line Valid [11]

Since the MT9V034 module transmits image data in parallel and each pixel contains 10 bits of resolution, 10 pins are used to transmit pixel values. Pixel data is transmitted in correspondence with LINE_VALID and output clock signal PIXCLK. When LINE_VALID is asserted, the pixel data pins are updated with values corresponding to pixels 0-751 of

the given line. Values for each pixel are written out on the falling edge of the camera's PIXCLK pin, allowing for each pixel's value to be read on each rising PIXCLK edge. A full LINE_VALID data transmission sequence will therefore contain 752 PIXCLK cycles, corresponding to the 752 pixels that make up the given line. A timing diagram of this data transmission scheme is shown in Figure 9.

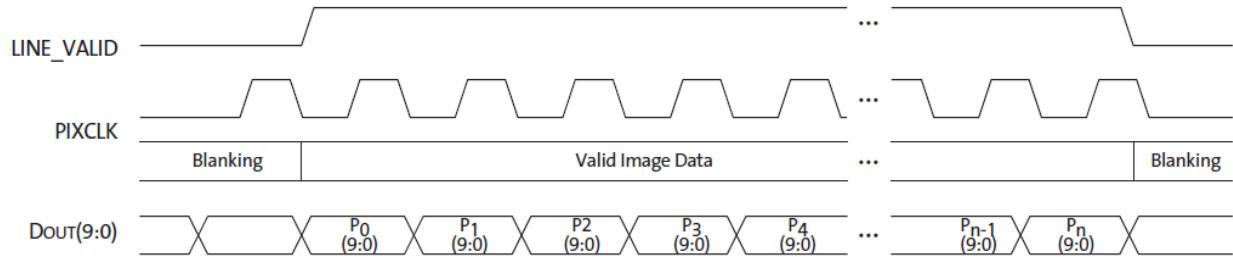


Figure 9: Line Data Transfer [11]

The default camera data transmission scheme was also examined using an oscilloscope, as shown in Figure 10, with channels 1-4 corresponding to camera PCLK, FRAME_VALID, LINE_VALID, and Data[0], respectively. In the case of Figure 10, the camera is initially powered off, resulting in an inactive PCLK signal during the beginning of the recording.

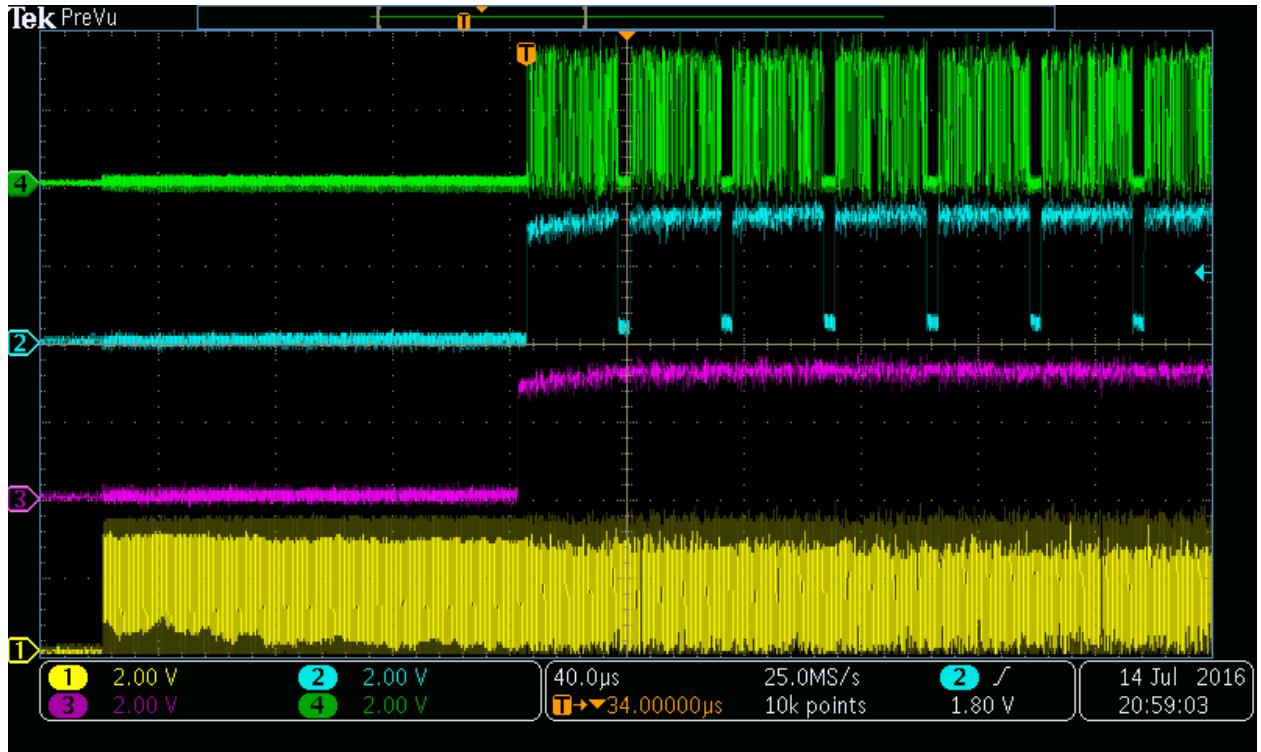


Figure 10: Camera Data Transfer

3.1.2 I²C Control

The MT9V034 Camera module's mode of operation can be configured using a standard I²C control interface. I²C, or Inter-Integrated Circuit, is a bidirectional serial interface that allows for a master device to read from and write to several slave devices sharing the same data bus. An I²C interface will use a Serial Data Line (SDA) and Serial Clock Line (SCL) that are normally pulled to 5V. When one connected I²C device wishes to communicate with another, it will pull the SDA line low while leaving the SCL line high. The master device will then begin clocking the SCL line, and SDA will be used to transfer 7 bits representing the address of the desired slave device, along with an 8th bit representing whether it would like to read from or write to the device. An example of this transfer is shown in Figure 11. A second 8 bit sequence representing a specific register within the slave device may also be transmitted following the device address. For example, if the master device wishes to write to slave device 0x40 at register 0x00, it will transmit 0x41 (address 0x40 and WRITE),

followed by 0x00. If the slave device receives this transmission, it will acknowledge by pulling the SDA line low. At this point, the master can then transmit the value that it wishes to write to the given slave address and register. If the operation were a read rather than a write, the slave would transmit a value back to the master.

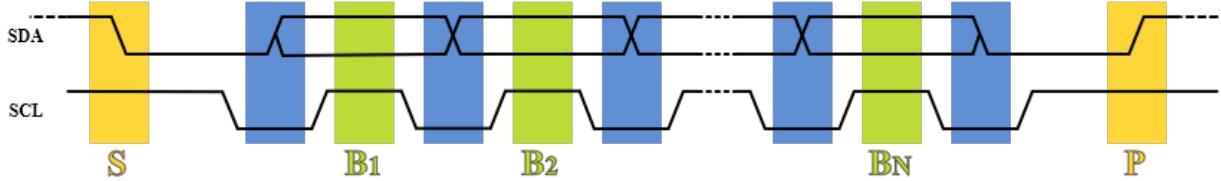


Figure 11: Example I²C Data Transfer

Based on the LIVM34LP camera board schematic, each breakout board have been configured so that its camera is accessible at I²C address 0x58 [5, 11]. Note that since both cameras come configured with the same I²C bus address, a pullup resistor must be added to one of the cameras I²C address lines so that both are individually accessible on a shared bus.

Although the MT9V034 camera control registers are closed source, the previous model's registers are available in the camera module datasheet, and have been found to work with the current model thus far [10]. As a baseline, the camera module was sent a read request at address 0x00, which should return 0x1324 for the MT9V034 camera module. An oscilloscope screenshot of this request is shown in Figure 12, with the first packet consisting of a request to address 0x00 of device 0x058, and the second packet consisting of the camera's response of 0x1324.

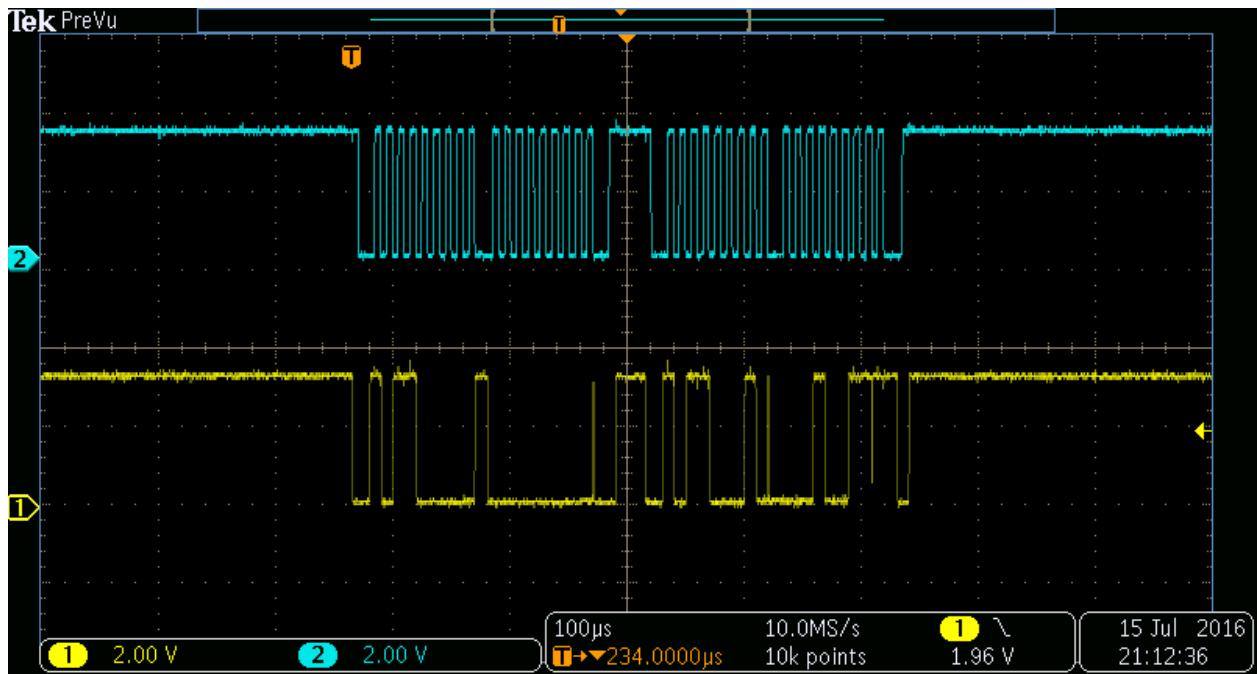


Figure 12: Example I²C Transfer with Camera

After the camera I²C was deemed working, the camera control register was modified to put the camera in “snapshot” mode. In this mode, the camera module will no longer continuously take pictures, and will only gather new images when an external trigger is activated. This is the mode that each camera will need to operate in in order to acquire stereo imagery, since a shared trigger line will allow for both cameras to be controlled simultaneously.

According to the previous camera iteration’s datasheet, the camera module’s operational mode can be set through control register 0x07. By default, this register will be set to a value of 0x0388, which corresponds to master mode with parallel output and simultaneous readout of pixel data enabled [10]. In order to put the camera in trigger mode, the control register needs to be written with value 0x0198, which allows for the same functionality as before with the exception of having continuous shutter mode replaced with an external trigger. For reference, a table with bit descriptions for the camera control register can be found in Appendix item 4.4 [10].

After modifying the state of this register, a button was attached to the camera’s TRIGGER input line, and the TRIGGER and FRAME_VALID lines were observed on channels one

and two of the oscilloscope, as shown in Figure 13. This oscilloscope screenshot can be seen as an example of how the camera is no longer in continuous operation, since FRAME_VALID only asserts itself in response to a TRIGGER input.

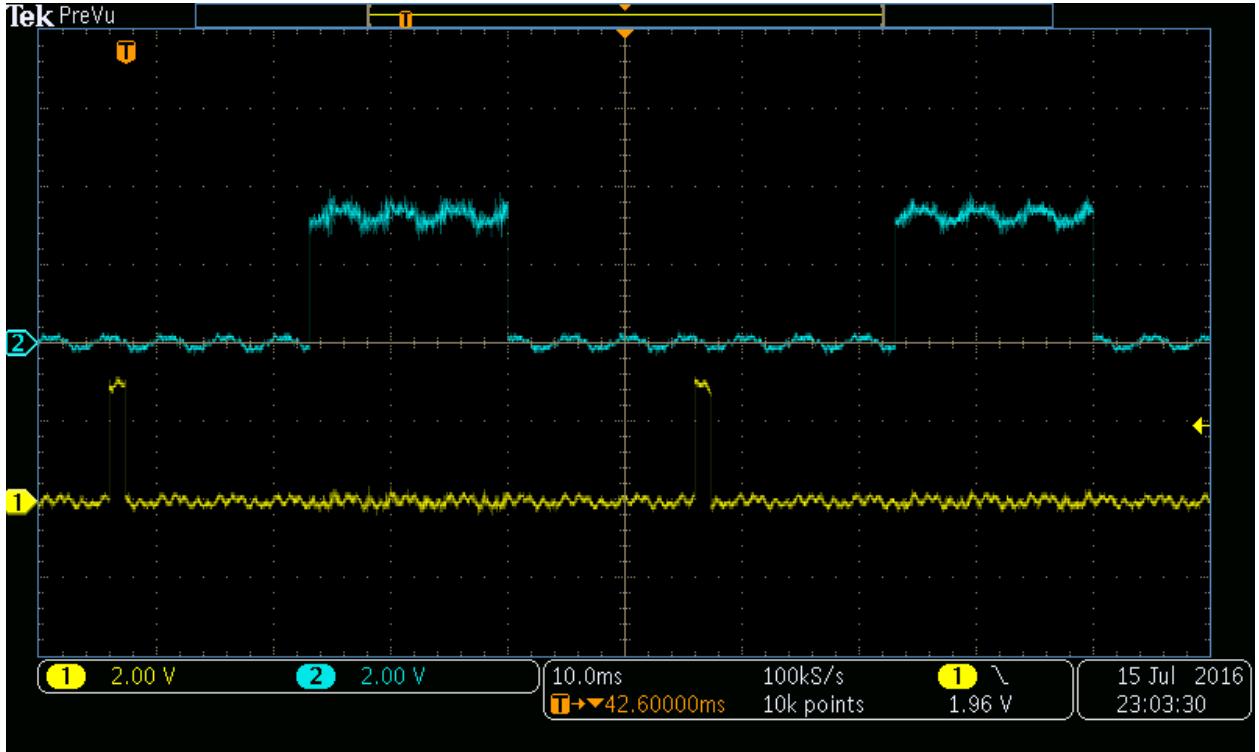


Figure 13: Camera Trigger and FV in Trigger Mode

In order to prevent accidental modification of the camera module's configuration registers, the register lock feature of the camera I²C bus is also used. By writing 0xDEAD to register 0xFE, it is possible to disable the I²C bus from being written to. This feature is disabled when the power of the camera module is cycled, or by writing 0xBEEF to the register lock register.

3.1.3 Data Management

After successfully creating a camera control interface and placing the MT9V034 camera module in trigger mode, it was then possible to begin viewing images from the module. Since each image contains 752x480 pixels with 10 bits of resolution per pixel, a full camera image

will consume 3,609,600 bits, or 440.6kB, as shown in Equation 1.

$$\text{Image Size} = 752\text{px} * 480\text{px} * 10 \frac{\text{bits}}{\text{pixel}} = 3609600 \text{ bits} * \frac{1 \text{ byte}}{8 \text{ bits}} * \frac{1 \text{ kB}}{1024 \text{ bytes}} = 440.625 \text{ kB} \quad (1)$$

In order to send a camera image to a computer or monitor for viewing, several steps need to be taken. Although it would be ideal to transfer the image directly from the camera to a computer or display, this would be difficult to achieve due to the high speeds of the camera's data output. In order to properly synchronize camera data with a VGA display, both the camera and VGA display would have to run at exactly the same clock speed, and would need to have the same amount of vertical and horizontal blanking to display each pixel in its correct location. If the image were transferred to a computer, the act of packaging the information so that it may be interpreted by said computer would place severe limitations on the speed of the system. A proper solution to these timing issues would be to buffer the image between the camera and the desired output source, since this would allow for separate clock domains to be used for camera data transfer and data output. However, the act of locally buffering a camera image on an FPGA would also be difficult due to low memory resources.

Although 440kB may seem like a relatively small image size, creating a buffer object large enough for storing said image would consume an extremely large amount of logic. For reference, a standard Nexys3 FPGA evaluation board contains only 18kB of onboard Block RAM (internal memory), and would not be able to buffer an image of this size without the use of external memory¹. This leaves the final option of using either external memory or a First-In First-Out (FIFO) memory array for transferring a captured image between clock domains. During initial development, an AL422B FIFO IC was used, since the IC has been created specifically for buffering VGA imagery similar to that of the MT9V034 camera module, and can be connected directly to the camera module outputs [1]. The AL422B

¹Xilinx, *Spartan-6 FPGA Block RAM Resources*, 11.
http://www.xilinx.com/support/documentation/user_guides/ug383.pdf

FIFO module contains 3M-bits of RAM that can be written to and read from in parallel, and supports separate input and output clock speeds between 1-50MHz [1]. This means that the camera module can write pixel data to the FIFO as long as it operates at a speed between 1 and 50MHz, and the FPGA can independently read from the FIFO at any speed within the same range. Note that since this FIFO supports only 8-bit parallel data in and out, the lowest two bits of camera pixel data must be truncated. This isn't a major issue, since the truncation will correspond to a 4/1024 reduction in the range of values that each pixel can map to.

With the inclusion of the external FIFO module, it is now possible to capture and store an image for future reading, and to read out image data in chunks. Keeping this in mind, the system shown in Figure 14 was created for capturing, storing, and transmitting camera images to a computer for external analysis. In order to reduce development time, an external microcontroller was used for controlling the camera module's I²C interface and placing the module in trigger mode. Various buttons and switches on the FPGA were then used for controlling the camera output and trigger, allowing for a user to trigger an image for storage on the AL422B FIFO. Once the image has been stored on the FIFO, the FPGA is capable of reading the image line-by-line into an internal buffer. An internal System on Chip (SoC) is used to control FPGA reads from the FIFO into this internal buffer. An image dump will begin when the SoC microcontroller signals to the FPGA to read a new line of pixels into its internal 8 bit by 752 address pixel buffer. The FPGA will then signal to the microcontroller when this buffer has been filled, and the microcontroller will print out the value of each pixel in the buffer to a connected computer over a Universal Asynchronous Reciever/Transmitter (UART) port. When the microcontroller finishes printing out the value of each pixel in the line buffer, it will signal to the FPGA to read in a new line of pixels. This process will repeat for each of the 480 lines of pixels in the image, allowing for the transmission of an entire image's worth of data from FIFO to computer. The Verilog implementation of the top module and line buffer for this interface can be found in Appendix item 4.7.1.

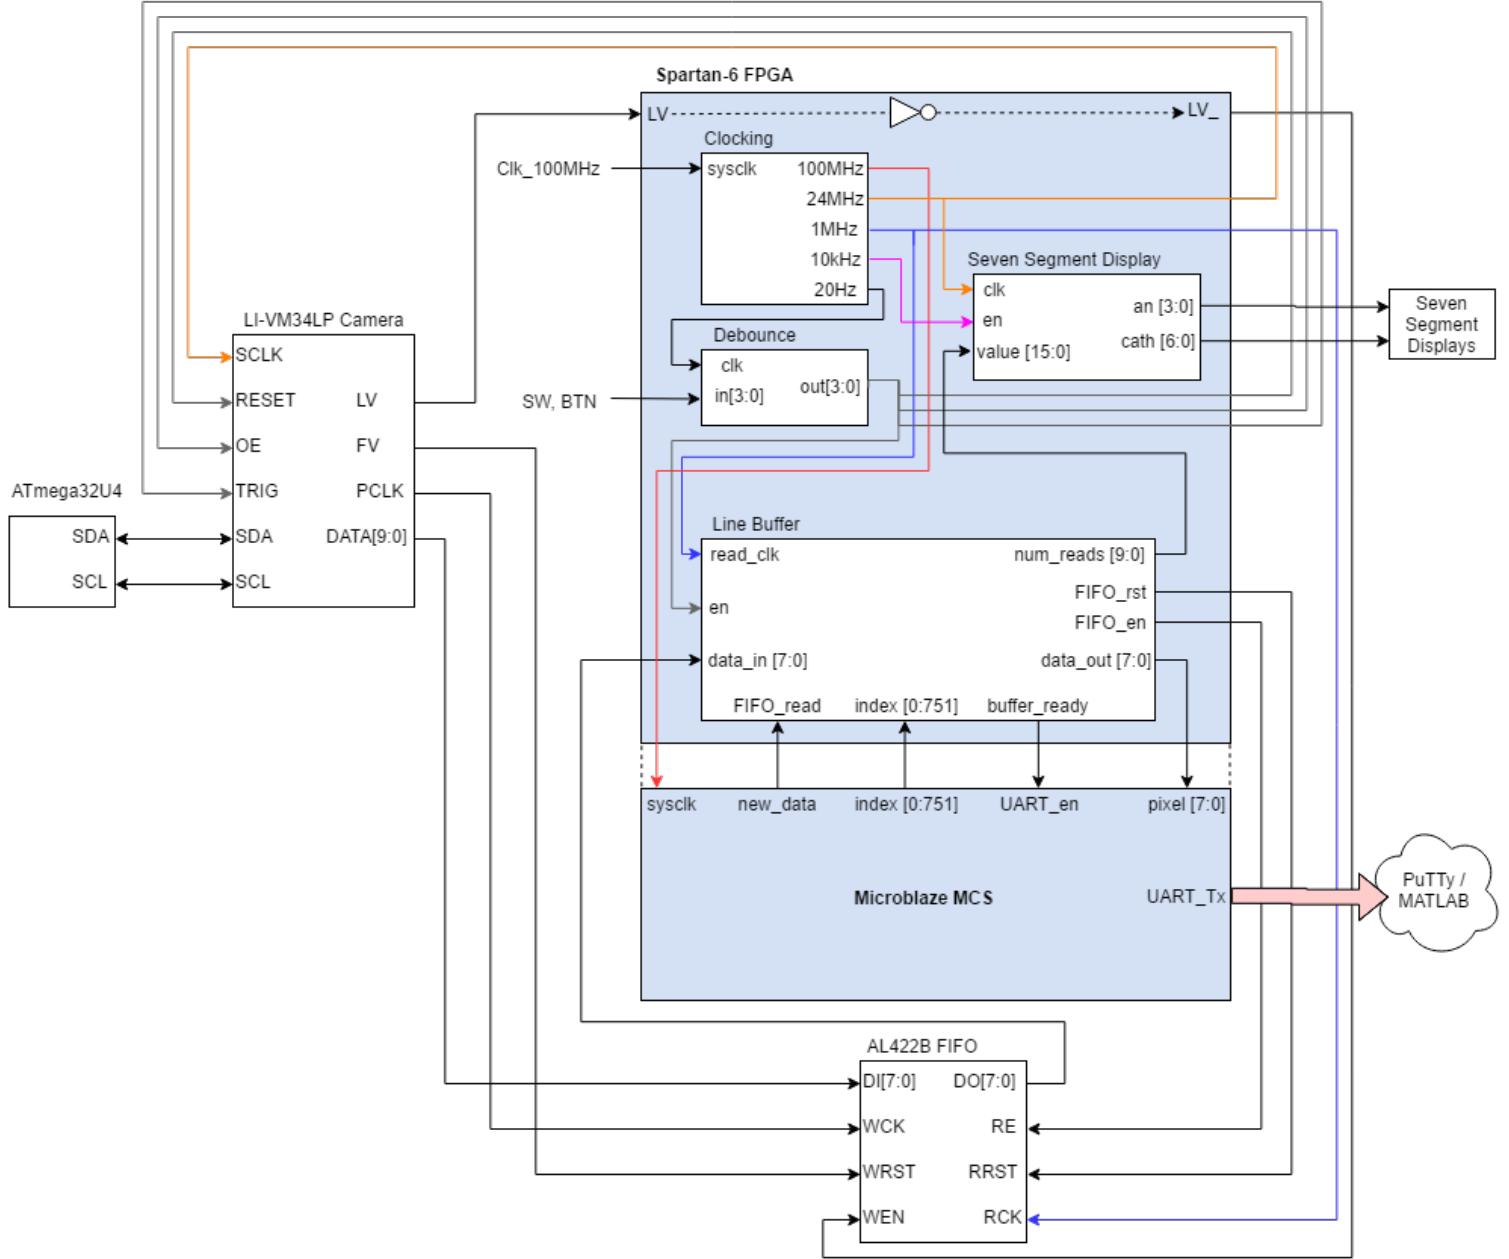


Figure 14: Camera Test System Block Diagram

An example of the transmission of one line of pixel data from the FIFO to the FPGA is shown in Figure 15. The green, purple, blue, and yellow lines in this image represent pixel data, FIFO read enable, read reset, and read clock, respectively. Since the FPGA reads in one line of pixel data at a time, this process will take 752 read clock cycles, as measured in Figure 15. In order to simplify debugging, an internal counter and seven-segment display

controller have also been implemented on the FPGA, and will display a running count of the number of pixel lines that have been read into the FPGA's internal buffer, ranging from 0x0000-0x01E0 (0-480).

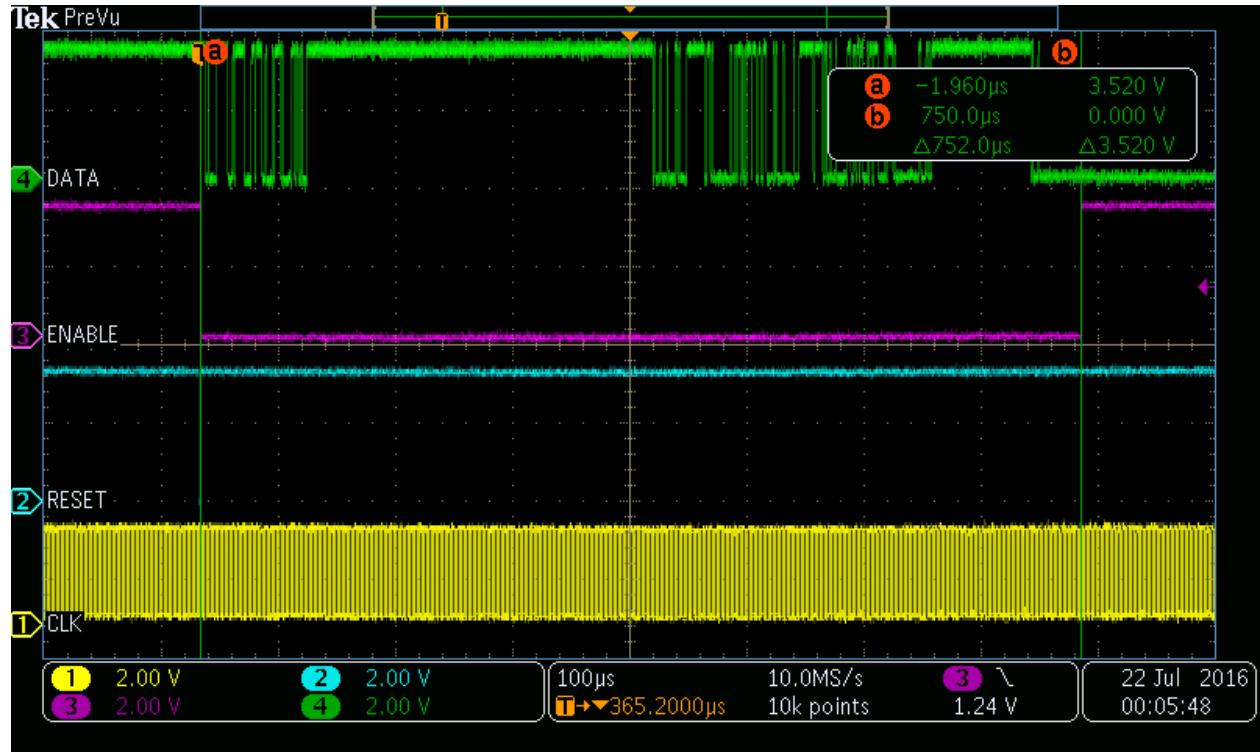
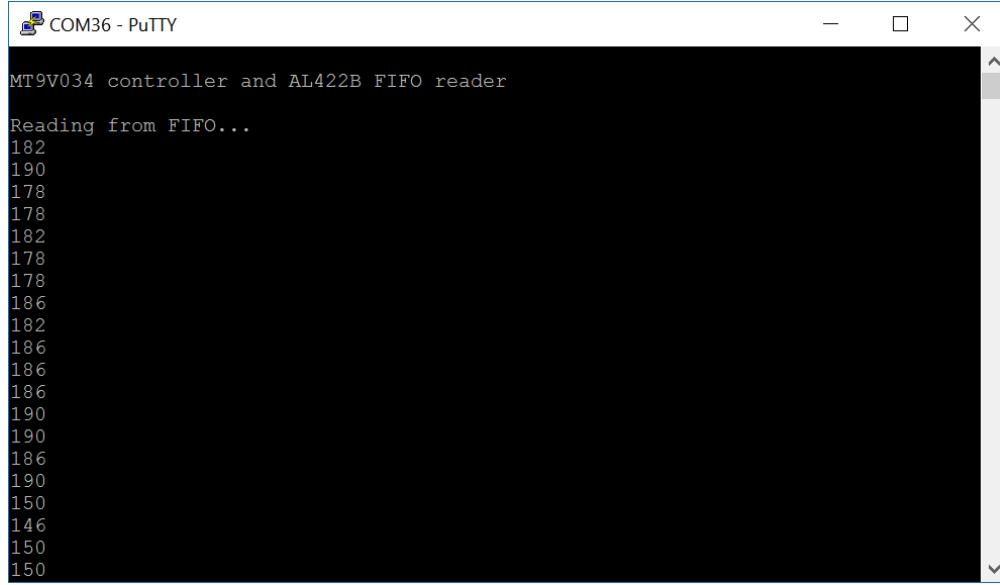


Figure 15: Transferring Line Data from FIFO to FPGA

3.1.4 Transmitting Images Over UART for Analysis

Once the FIFO and FPGA line buffer interfaces were created, the source code found in Appendix item 4.8.1 was implemented on a Microblaze SoC in order to transmit camera line data from the FPGA's internal line buffer over UART. An example of the microcontroller's UART output is shown in Figure 16. The microcontroller will print the value of each pixel followed by a newline and carriage return, starting with the top left pixel in the acquired image.



```
MT9V034 controller and AL422B FIFO reader
Reading from FIFO...
182
190
178
178
182
178
186
182
186
186
186
190
190
186
190
150
146
150
150
```

Figure 16: Reading FIFO Data

After the image is received through PuTTY, the MATLAB script found in Appendix item 4.6.1 is used to parse the corresponding logfile into a greyscale image. An example image created through this process is shown in Figure 17. Note that the sub-optimal quality of this image is due to signal interference and degradation in the test setup's long wiring, as shown in Figure 18.



Figure 17: Notebook With Grid and Oscilloscope Leads

Although this system was tested using the Nexys3 (Spartan-6) FPGA board, the use of an external FIFO and little to no platform-specific hardware make it so that it can easily be implemented on any system, including the Zynq family of processors that we will be implementing our final system on.

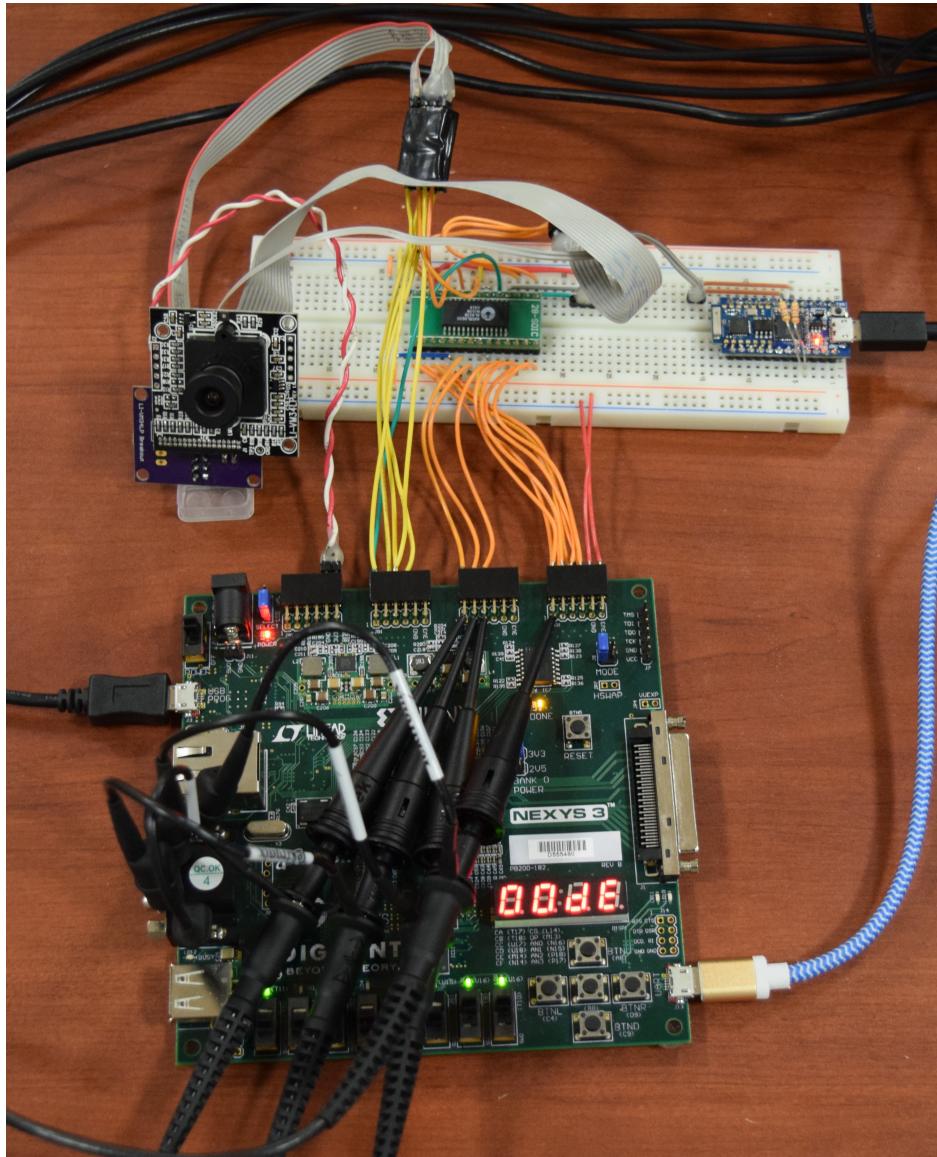


Figure 18: Camera Test Setup

3.2 Final Camera Hardware Implementation

After successfully gathering image data from a single camera module, an interface needed to be created for controlling both cameras at once using the ZedBoard. Interfacing each camera module directly to the ZedBoard's GPIO is not feasible, since the pair would consume every available PMOD pin on the board, leaving no additional pins for the IMU or rangefinder². One solution originally investigated was the use of the ZedBoard's FPGA Mezzanine Card (FMC) connector, since it contains 68 available GPIO pins and would be more than adequate for interfacing the stereo cameras with the board. However, the FMC connector has been configured to provide logic voltage levels of only 1.8 or 2.5 volts without modification to the ZedBoard. Since each camera module is only compatible with 3.3 volt logic, the FMC connector is therefore not be feasible for our designs.

This leaves the final option of reducing the overall pin count required by the cameras and interfacing the combined camera setup with the board's PMOD pins. One significant method of reducing the necessary pins required is to include an individual AL422B FIFO per camera. Based on the testing described in the previous section, it has already been determined that these FIFO modules are compatible with the MT9V034 cameras, and are capable of significantly reducing memory requirements on the FPGA. A second major advantage of including these FIFO modules in the camera interface is that their data output lines may be placed in a high-impedance state. This means that the individual data output lines of each FIFO module can be connected in parallel, with a single FIFO driving the lines at a time. Since the bulk of each camera module's required pin count lies in its data lines, the ability to connect these lines in parallel reduces the overall camera GPIO requirements by 8 pins. Since each AL422B FIFO module is capable of being read from at a clock speed of up to 50MHz and the maximum master clock rate of each MT9V034 camera module is 27MHz, the inclusion of the FIFO modules also won't cause a significant decrease in the overall speed of the stereo camera system [1, 11].

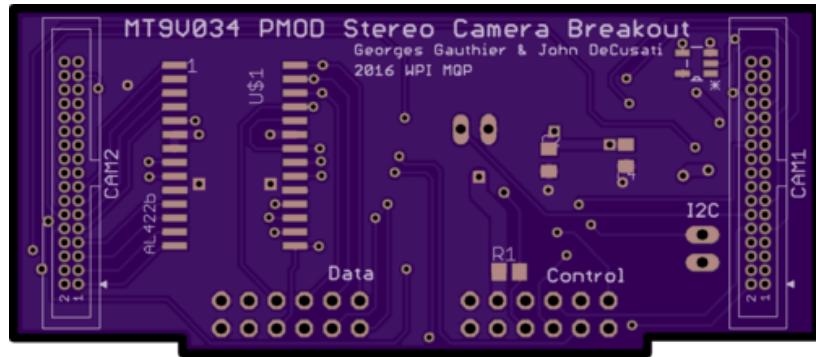
² $[2 * (D[9 : 0] + TRIGGER + OE + RST + SCLK + PCLK + FV + LV)] + SDA + SCL = 36 \text{ pins}$

Along with the shared camera data lines between each AL422B FIFO module, it is also possible to connect several other signals in parallel. Since each camera image capture should be triggered at approximately the same time in a stereo imaging setup, it is already desirable to connect both camera TRIGGER lines together. The RST, OE, SDL, SCA, and SCLK lines of each camera module can also be tied together in pairs of two, and the OE lines can simply be held at 3.3 volts. Lastly, since each camera LV signal must be inverted for use with the AL422B FIFOs, a discrete inverter IC may be used to save on FPGA GPIO. Overall, these modifications will save a total of 25 pins, as shown in Equation 2.

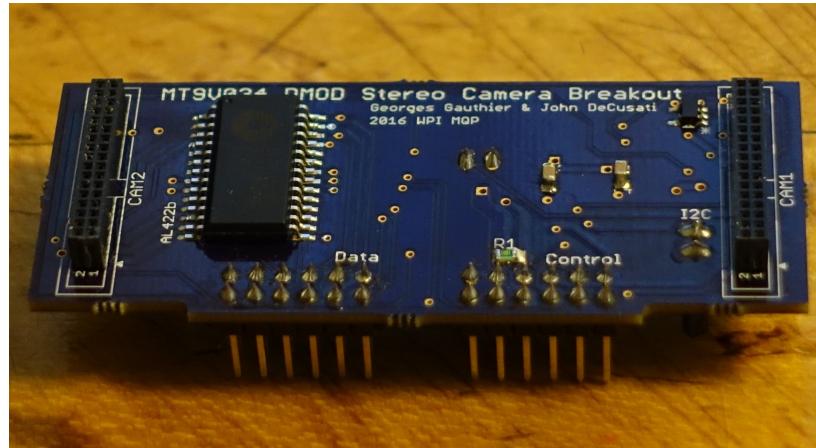
$$36 \text{ Pins} - (8 \text{ Data} + 4 \text{ truncated bits}) - (\text{TRIGGER} + \text{SCLK} + \text{RST}) \\ - 2 * (\text{OE} + \text{PCLK} + \text{FV} + \text{LV}) = 13 \text{ pins (!)} \quad (2)$$

Note each FIFO must be controlled individually, requiring an additional Read Reset (RRST) and Read Enable (RE) pin per camera, as well as a shared Read Clock (RCK) line. This brings the total pin count required by the stereo camera setup to 16 pins plus two I2C pins, which is conveniently the number of GPIO available in two PMOD headers. This setup was implemented as shown in Appendix item 4.5, and the final stereo camera breakout board shown in Figure 19 was then created.

A Verilog module was created using a modified version of the MT9V034 camera test code found in Appendix item 4.7.1 and a VGA controller in order to test the stereo camera breakout board. A switch input is used to select one of the two camera modules for image acquisition, and a binned 60x92 pixel set from the center of the camera's image is buffered locally for VGA display. The image is then independently written to the display according to internal VGA timing. This process is repeated at a high rate of speed, allowing for a realtime video stream from the selected camera to be displayed. The assembled stereo breakout board used in this test is shown in Figure 20.



(a) PCB Top



(b) Assembled PCB

Figure 19: Stereo Camera PMOD PCB

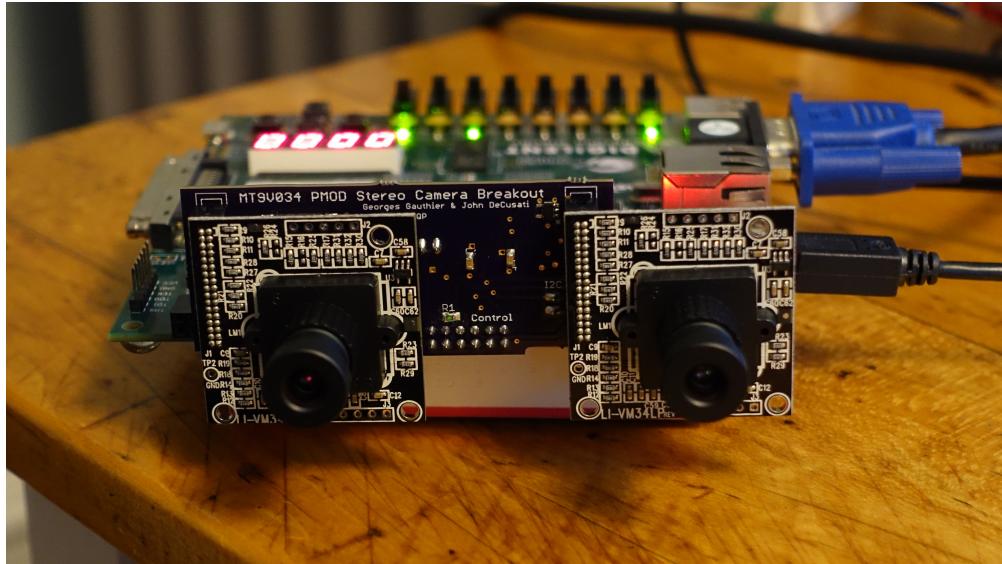


Figure 20: Stereo Camera Breakout Under Test

After attempting to manually focus each camera using the VGA module described above, the code used in Section 3.1.4 was used to transmit image data from the stereo cameras to a computer for further analysis. As you can see from the example image in Figure 21 below, the new stereo camera setup is far less susceptible to data loss in comparison to the previous version.

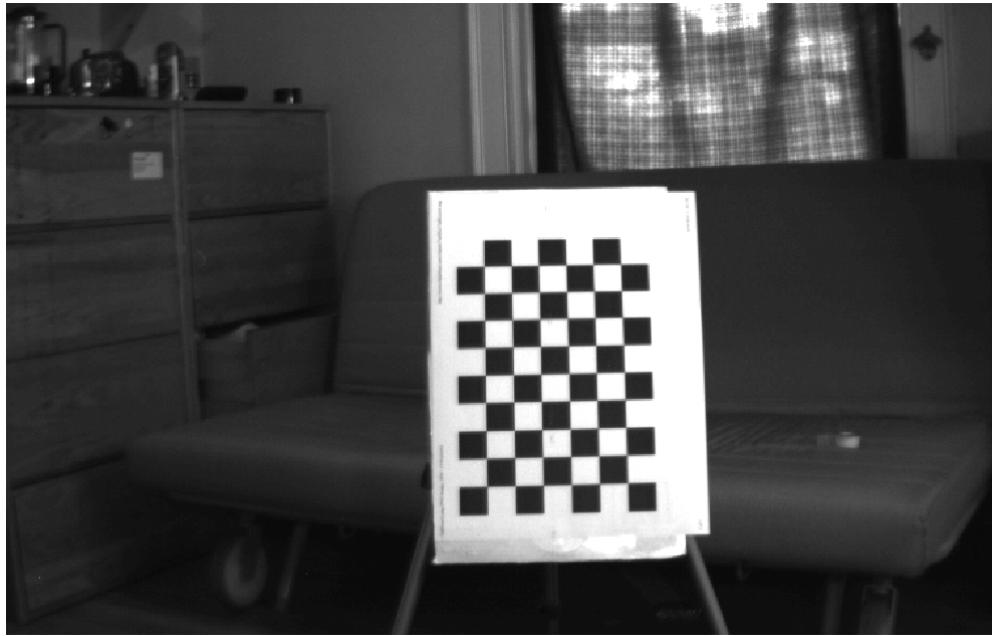


Figure 21: Stereo Camera Breakout Sample Image

For further comparison, please refer back to the test image acquired using the original camera test setup, as shown in Figure 17.

3.3 Disparity Algorithm

After verifying that the camera interface was functional, a large portion of time was spent implementing a disparity algorithm that would allow for the extraction of 3D depth information from stereo image data. This algorithm was first implemented in MATLAB, and was then transferred to programmable logic after the algorithm was verified working.

3.3.1 Sum of Absolute Differences

The method used in our disparity algorithm implementation is known as the Sum of Absolute Differences, or SAD. SAD is a common digital image processing technique used to measure the similarity between blocks of image data. In the case of our stereo camera interface, a SAD algorithm is used to search through the right image data for windowed blocks that match a template block selected from the left camera image. This process is performed using 7x7 pixel search blocks over 50 pixel horizontal ranges, and is repeated throughout the image. The expression for the sum of absolute differences is shown in Equation 3 below.

$$SAD = \sum_x \sum_y |template - block| \quad (3)$$

A visual representation of the Sum of Absolute differences is shown in Figure 22, with the top image showing the left image template block, and the middle image showing the right image search window in relation to the location of the template block. Below both images is a visual representation of the Sum of Absolute Differences between the template block and the current search block, outlined in white. In the case of the current search, the template and search blocks are relatively different, resulting in a high SAD value.

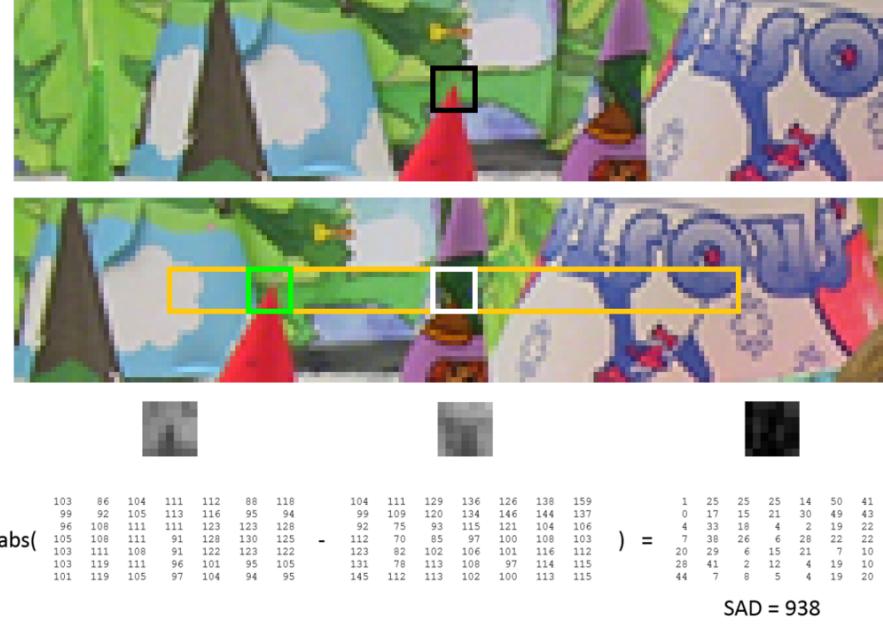


Figure 22: Sum of Absolute Differences [9]

Since the disparity algorithm used in this implementation calculates the sum of absolute differences for multiple search blocks, the resulting SAD values for each search block can be compared to find the location of the most similar matching block in the search image. Due to the nature of the SAD algorithm, lower SAD values indicate higher similarity between the template and search blocks. This comparison is demonstrated in Figure 23 below. In the case of Figure refblockMatching, higher match score values for each search block indicate lower SAD values.

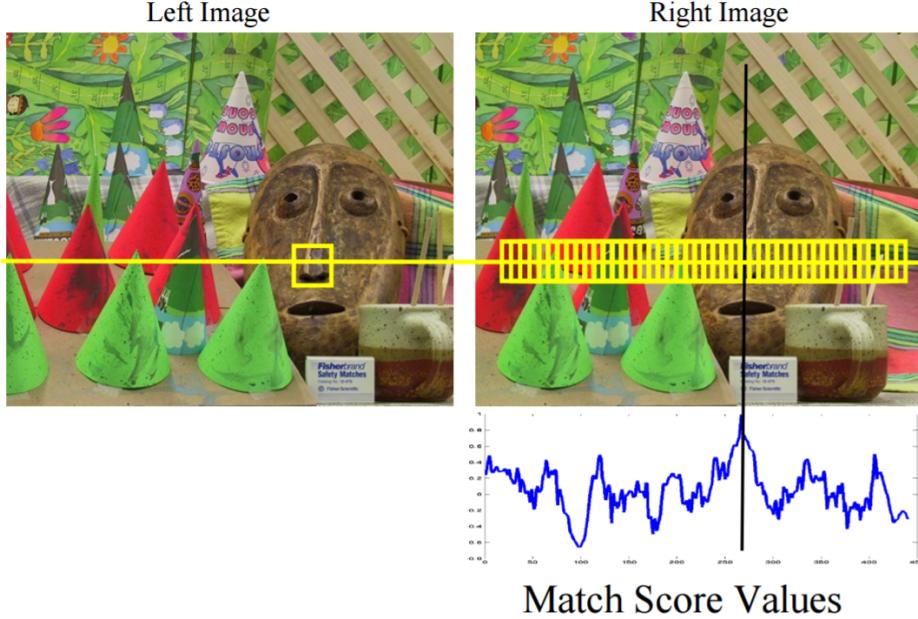


Figure 23: Block Matching Overview [2]

The SAD at multiple search points can be used to estimate the pixel offset between the template block and matching search block based on array index locations, since all SAD values for a single search are stored in a vector. This pixel offset is known as the disparity value for a given template and search block. The disparity d at a given point can be transformed into units of distance using the focal point f and baseline distance T_x between image sensors as shown in Equation 4 below.

$$\text{depth} = Z = \frac{fT_x}{d} \quad (4)$$

Pixel coloration values in a disparity image are based on the distance calculation shown in Equation 4, where each pixel is referenced to the disparity at a given template block's location. An example disparity image created using the MATLAB code found in Appendix item 4.6.2 is shown in Figure 24 below.



Figure 24: Disparity Algorithm Output

3.3.2 Image Rectification

The Sum of Absolute Differences algorithm operates under the assumption that objects in both camera images lie on the same horizontal line between both images, known as an epipolar line. An example of shared epipolar lines between camera imagery is shown in Figure 25 below. Although an ideal stereo camera setup would contain shared epipolar lines between camera images, raw image data from each camera will contain slight differences in object location based on the physical position of the camera modules, as well as minor differences in the lenses of each camera. Both input images can be adjusted to share the same epipolar lines through a post-processing step known as image rectification.

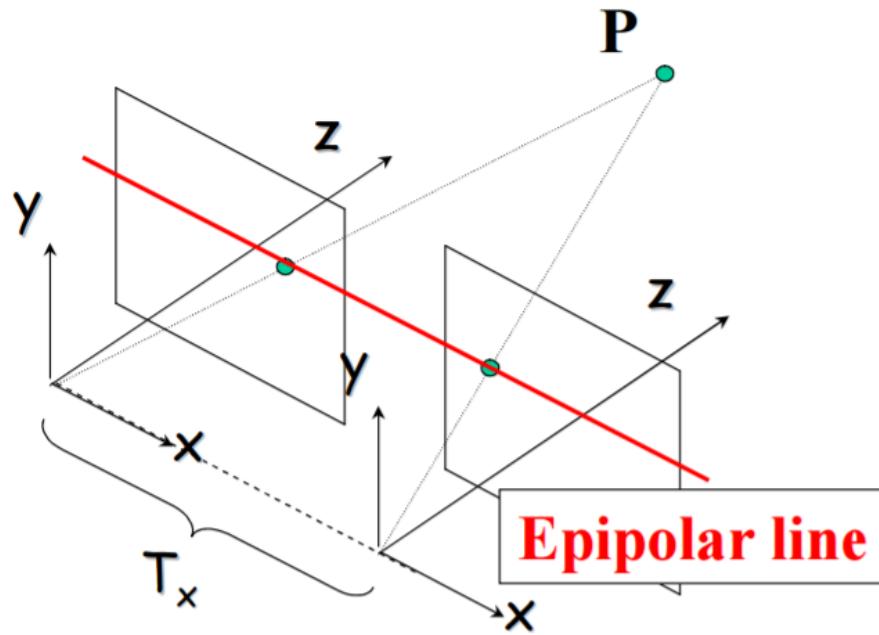


Figure 25: Horizontal Epipolar Lines [2]

A pictorial representation of the process of stereo image rectification is shown in Figure 26 below [8]. This process is achieved using a 3×3 matrix coordinate transform based on parameters obtained from the external calibration process.

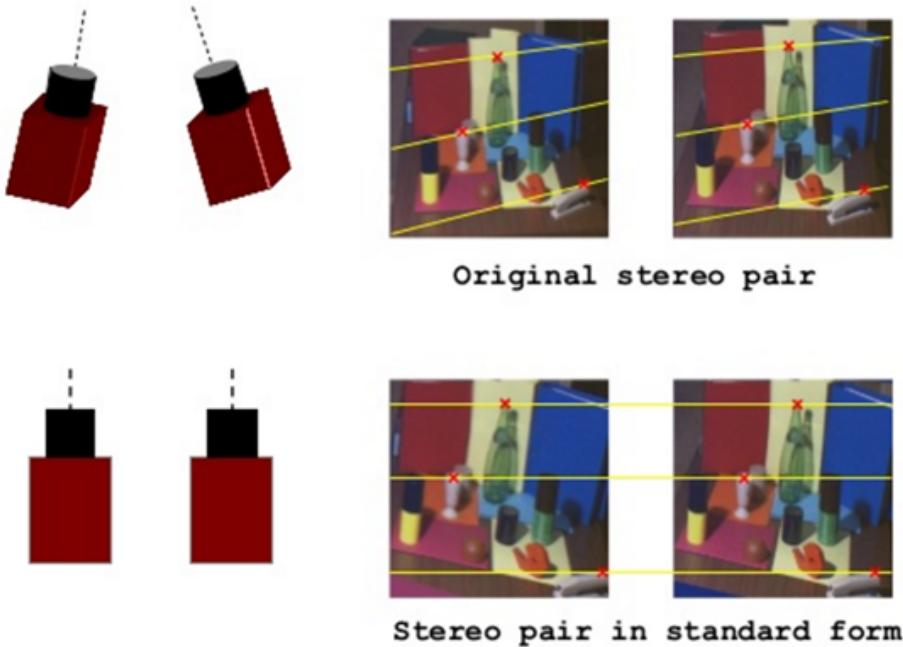


Figure 26: Stereo Image Rectification [8]

ADD AN EXAMPLE OF THE MATLAB CALIBRATION STUFF (figures showing 3d anaglyph process)

3.3.3 Test Implementation

The original disparity test implementation used closely follows the MATLAB disparity algorithm shown in Appendix item 4.6.2. This algorithm is implemented using a finite state machine with five states, as shown in Figure 28 below. In order to maintain simplicity, the test algorithm has been implemented to operate on 46x30 windowed portions of the input imagery. By default, the disparity module will remain in an idle state until an external enable signal is toggled high using a button input. This will cause the finite state machine to advance to its READ state, and image data for the left and right camera images will be read in from the stereo camera breakout board. After image data has been received, the state machine will then advance to a cyclical set of states used for iterating through each image and calculating disparity.

The disparity module will begin by isolating the template and search blocks from the

right and left image data in the finite state machine's separation state. Next, the state machine will advance to its SAD state, and will calculate the sum of absolute differences between the template and search block. This value is placed in a vector that matches the length of the search range. If the vector hasn't been completely filled, indicating that there are more search blocks to compare to the template, the state machine will revert back to the separate state, isolating a new search block from the right camera image. When the SAD vector is full, the state machine will advance to its finalization state.

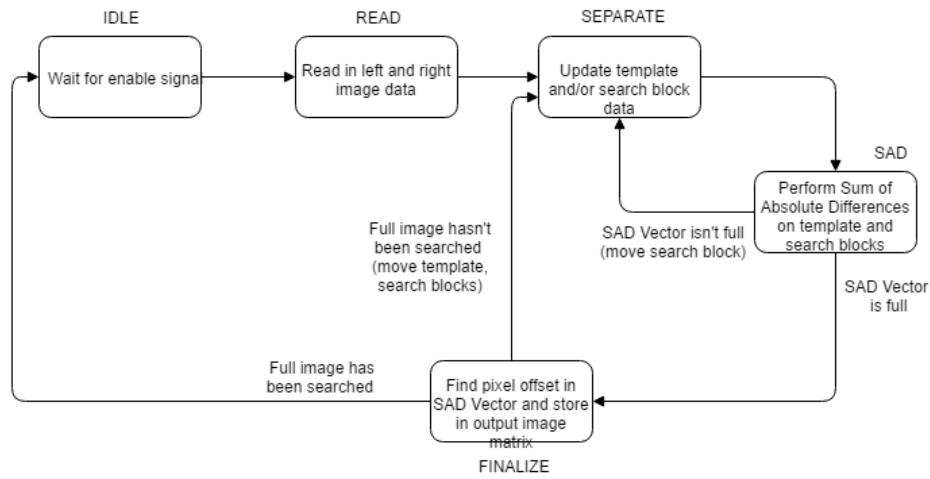


Figure 27: Disparity Test Implementation

The finalization state is used to search through the SAD vector for the lowest value. The index of this value within the SAD vector in reference to the template block location is used to create a disparity value for the given template block location. This value is then converted to a distance using Equation 4, and is stored in the output image location. If the output image hasn't been fully populated with distance values, the state machine will then revert back to the separate state. Otherwise, the state machine will advance to its idle state, and the resulting disparity image can be read for output.

This module was initially tested using a verilog test bench, and was then tested using camera image data and a VGA display controller module, allowing for real-time verification of the algorithm's effectiveness. After testing the initial disparity algorithm, several

modifications were made to increase the overall speed and efficiency of the disparity module.

3.3.4 Final Implementation

ZedBoard Disparity Algorithm Implementation

Zynq-7000 SoC + Onboard Memory

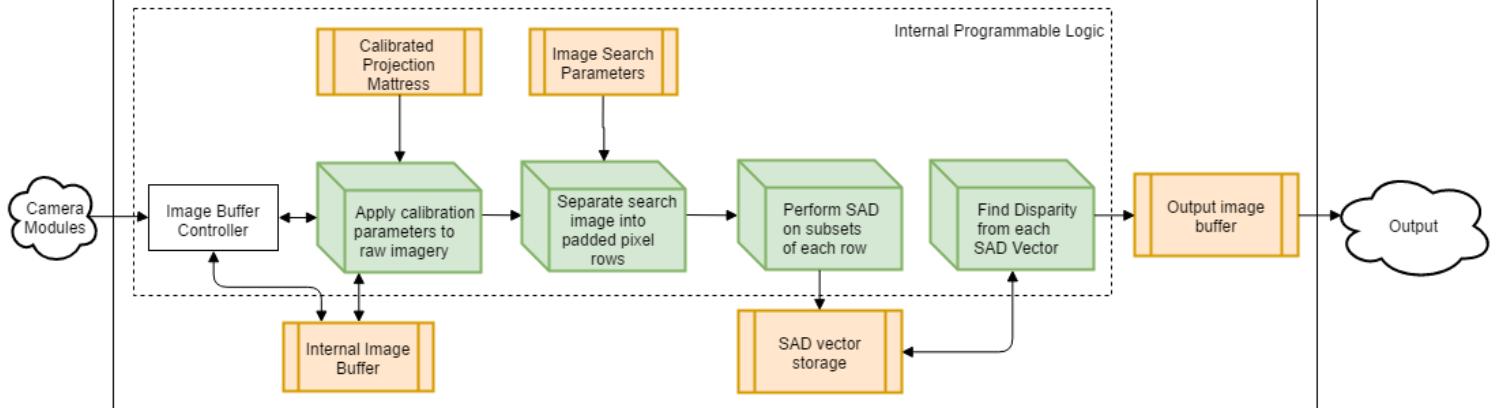


Figure 28: Disparity Final Implementation

References

- [1] Averlogic. *AL422B Data Sheets*. 2001. Available from: http://www.frc.ri.cmu.edu/projects/buzzard/mve/HWSpecs-1/Documentation/AL422B_Data_Sheets.pdf.
- [2] Robert Collins, *CSE/EE486 Computer Vision I Fall 2007 Lecture Notes*. Available from: <http://www.cse.psu.edu/~rtc12/CSE486/>.
- [3] Davison, A.J., *Real-Time Simultaneous Localisation and Mapping with a Single Camera*. IEEE Computer Vision, 2003. 2(1).
- [4] Dresser, L.H., A.W., *Accelerating Augmented Reality Video Processing with FPGAs*. 2016. Available from: <https://www.wpi.edu/Pubs/E-project/Available/E-project-042716-172155/unrestricted/armqp-final-report.pdf>
- [5] Leopard Imaging Inc. *LI-VM34LP Camera Board*. 2009. Available from: http://www.leopardimaging.com/uploads/li-vm34lp_v1.1.pdf.
- [6] *Serveball*. Available from: <http://www.serveball.com/>.
- [7] Stefano Mattoccia, M.P., *A passive RGBD sensor for accurate and real-time depth sensing self-contained into an FPGA*. in *International Conference on Distributed Smart Cameras*. 2015.
- [8] Stefano Mattoccia, *Stereo Vision*. Available from: <http://www.slideshare.net/DngNguyễn43/stereo-vision-42147593>
- [9] Chris McCormick, *Stereo Vision Tutorial - Part 1*. Available from: <http://mccormickml.com/2014/01/10/stereo-vision-tutorial-part-i/>.
- [10] On Semiconductor. *MT9V032: 1/3-Inch Wide-VGA CMOS Digital Image Sensor*. 2015. Available from: http://www.onsemi.com/pub_link/Collateral/MT9V032-D.PDF.
- [11] On Semiconductor. *MT9V034: 1/3-Inch Wide-VGA CMOS Digital Image Sensor*. 2015. Available from: http://www.onsemi.com/pub_link/Collateral/MT9V034-D.PDF.
- [12] Fatih Porikli, O.T., *Human Body Tracking by Adaptive Background Models and Mean-Shift Analysis*. in *IEEE International Workshop on Performance Evaluation of Tracking and Surveillance*. 2003.
- [13] Sebastian Thrun, D.H., David Ferguson, Michael Montemerlo, Rudolph Triebel, and C.B. Wolfram Burgard, Zachary Omohundro, Scott Thayer, William Whittaker. *A System for Volumetric Robotic Mapping of Abandoned Mines*. in *IEEE International Conference on Robotics and Automation*. 2003.

4 Appendix

4.1 Useful Resources

This section is intended to serve as complete compilation of all resources gathered throughout D Term 2016 that we believe will be useful as we begin to work on the methodology portion of our project.

- Strother, Daniel. 2011. “Open-Source FPGA Stereo Vision Core Released.” <https://danstrother.com/2011/06/10/fpga-stereo-vision-core-released/>.
- Field, Mike. 2013. “Zedboard OV7670.” http://hamsterworks.co.nz/mediawiki/index.php/Zedboard_OV7670.
- “OV7670/OV7671 CMOS VGA CameraChip Implementation Guide.” https://www.fer.unizg.hr/_download/repository/OV7670new.pdf.
- “MIPI CSI2-to-CMOS Parallel Sensor Bridge - Lattice Semiconductor.” http://www.latticesemi.com/~/media/LatticeSemi/Documents/ReferenceDesigns/JM/MIPICSI2toCMOSPar.pdf?document_id=50533.
- “OmniVision Serial Camera Control Bus (SCCB) Functional Specification.” http://www.ovt.com/download_document.php?type=document&DID=63.
- Morvan, Yannick. “Multiple-View Depth Estimation.” <http://www.epixea.com/research/multi-view-coding-thesisse15.html>.
- MathWorks. “Depth Estimation From Stereo Video.” <http://www.mathworks.com/help/vision/examples/depth-estimation-from-stereo-video.html>.
- Stefano Mattoccia, Matteo Poggi. 2015. “A passive RGBD sensor for accurate and real-time depth sensing self-contained into an FPGA”. International Conference on Distributed Smart Cameras.

- Mattoccia, Stefano. 2013. “Stereo Vision: Algorithms and Applications.” <http://www.slideshare.net/DngNguyn43/stereo-vision-42147593>.
- Szeliski, Richard. 2010. Computer Vision: Algorithms and Applications. New York: Springer.
- Beau Tippets, Dah Jye Lee, Kirt Lillywhite, James Archibald. 2013. “Review of Stereo Vision Algorithms and Their Suitability for Resource-Limited Systems.” Journal of Real-Time Image Processing 11 (1). doi: 10.1007/s11554-012-0313-2.
- Jouni Rantakokko, Joakim Rydell, Peter Stromback, Peter Handel, Jonas Callmer, David Tornqvist, Fredrik Gustafsson, Magnus Jobs, Mathias Gruden. 2011. “Accurate and Reliable Soldier and First Responder Indoor Positioning: Multisensor Systems and Cooperative Localization.” IEEE Wireless Communications 18 (2):10-18. doi: 10.1109/MWC.2011.5751291.
- Davison, Andrew J. 2003. “Real-Time Simultaneous Localisation and Mapping with a Single Camera.” IEEE Computer Vision 2 (1). doi: 10.1109/ICCV.2003.1238654.
- Fatih Porikli, Oncel Tuzel. 2003. “Human Body Tracking by Adaptive Background Models and Mean-Shift Analysis.” IEEE International Workshop on Performance Evaluation of Tracking and Surveillance.
- Giovanni Pintore, Enrico Gobbetti. “Effective mobile mapping of multi-room indoor structures.” The Visual Computer 30 (6):707-716. doi: 10.1007/s00371-014-0947-0.
- “iRobot 110 FirstLook.” iRobot. [http://www.irobot.com/\\$~\\$/media/Files/Robots/Defense/FirstLook/iRobot-110-FirstLook-Specs.pdf](http://www.irobot.com/$~$/media/Files/Robots/Defense/FirstLook/iRobot-110-FirstLook-Specs.pdf).

4.2 Component Selection

Component	Part Number	Supplier	Cost
FPGA	ZedBoard	Borrowed	N/A
IMU	ADIS16375	Borrowed	N/A
Rangefinder	URG-04LX	Borrowed	N/A
Stereo Cameras [†]	MT9V034	Mouser	\$146

[†] Note that we originally planned to purchase a flir lepton thermal camera module and accompanying breakout board to support two stereo ov7670 camera modules. After experimenting with the ov7670 camera module on our FPGA board, we began to realize that these camera modules are highly limited due to their low frame rate and poor documentation, and realized that we wanted to search for a different camera module. In addition, at a price of \$223 for a thermal camera with an 80x60 degree resolution, 25 degree fov, and 7-9Hz image sample rate, we believe that we are much better off spending our money on better camera modules that will be more usable for our task. For more information see Appendix Item 3.

4.3 Camera Module Decision Matrix

Camera Module	Max Frame Rate (FPS)	Resolution at Max Frame Rate (px.)	Cost	Requires External Adapter	Data Transfer Interface	Shutter	Field of View (deg.)	Rank 0-10
OV7670	30	640x480	\$10	No	Parallel	Rolling	25	5
Raspberry Pi Camera	90	640x480	\$30	Yes, \$53	MIPI (CSI2)	Rolling	49	6
PC1089K	60	720x480	\$32	No	NSTC/PAL	Rolling	Not Given	5
OV4682	330	640x480	\$89	Yes, \$50	MIPI	Rolling	Not Given	6
MT9V034	60	750x480	\$73	No	Parallel	Global	55	9

For purposes of comparison, the thermal camera module we were evaluating is also shown below.

Flir Lepton	9	80x60	\$175	Yes, \$40	SPI/ MIPI	N/A	50	3
-------------	---	-------	-------	-----------	-----------	-----	----	---

Shown above is our decision matrix for choosing a camera module. Fields marked in green indicate a positive ranking, while red indicates a negative ranking. Based on the individual

rankings of each item's field, we gave our camera modules an overall ranking of 0-10 in the right hand column, with 10 being an extremely high ranking and 0 being an extremely low ranking.

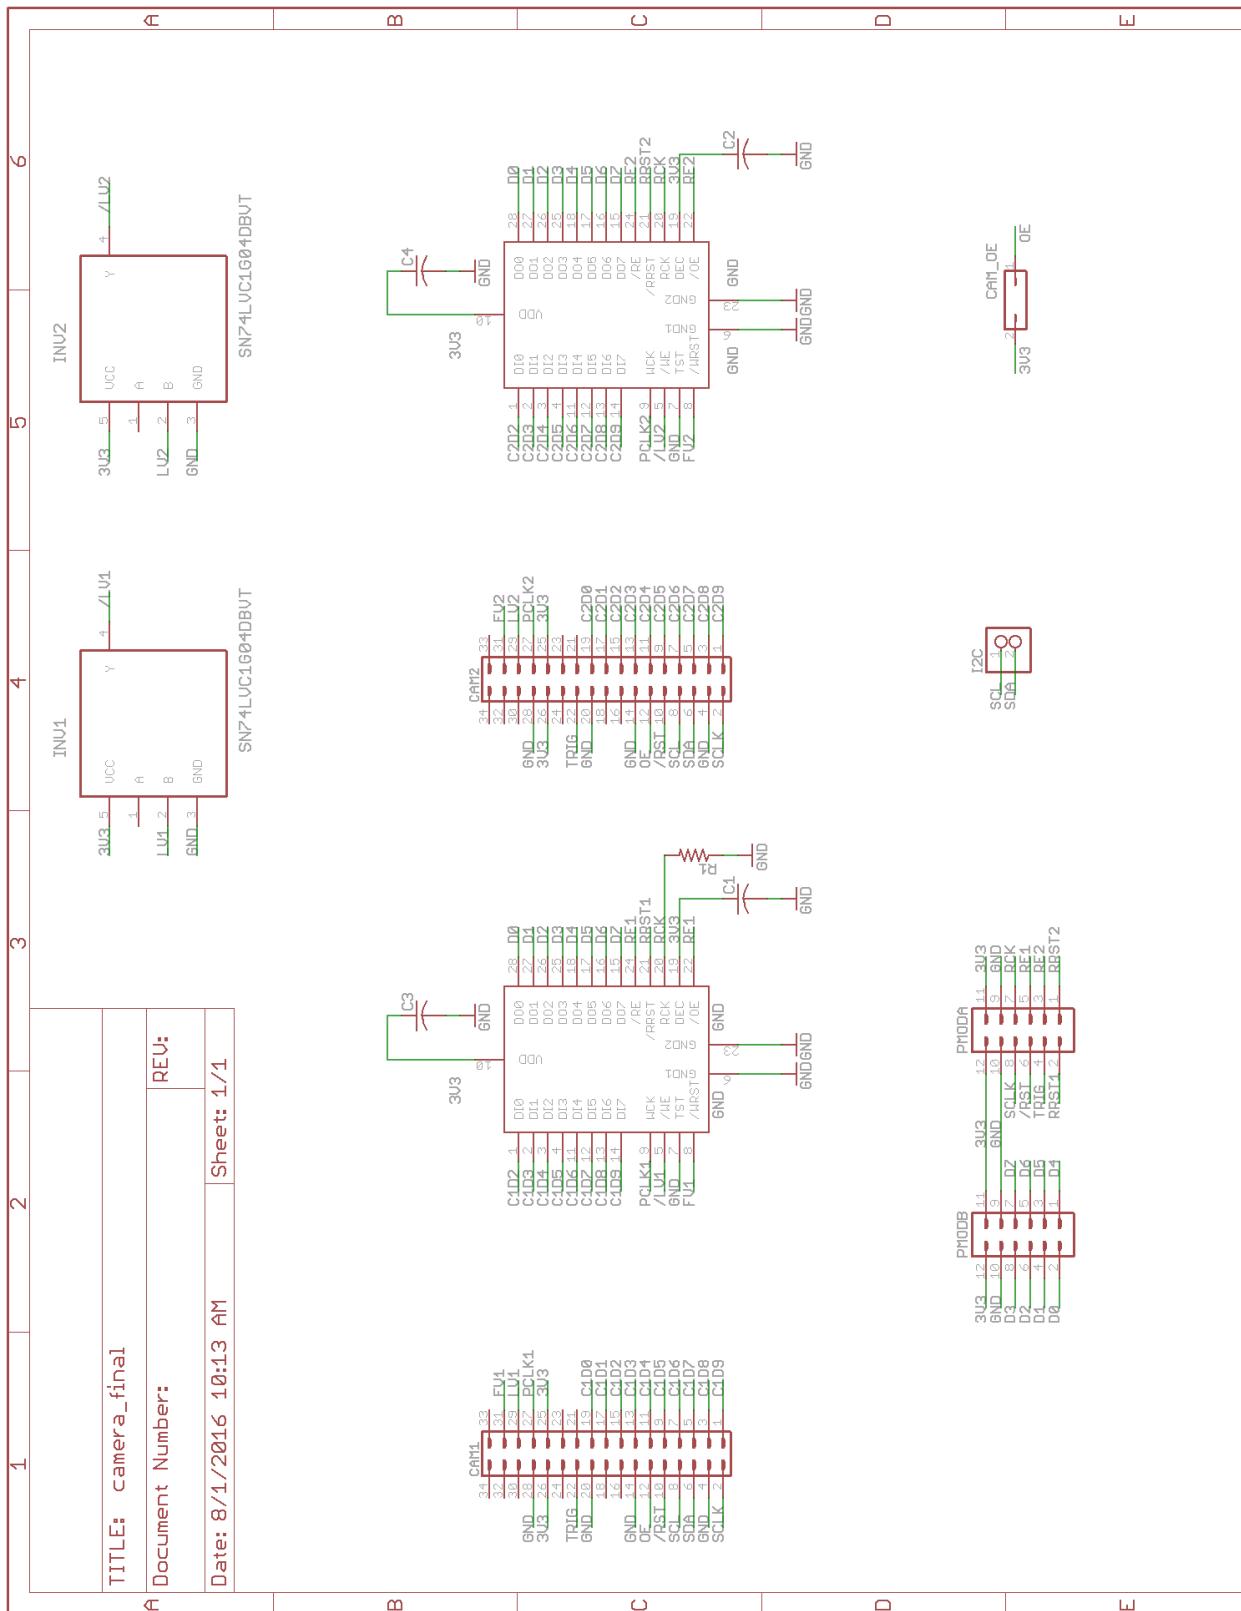
Based on our decision matrix, we believe that it would be worth both our time and money to use the MT9V034 camera modules for our stereo camera interface. These camera modules are the only low-cost global shutter option we've come across, and would be ideal for taking images in a sensor suite that is susceptible to motion. The MT9V034 also uses a parallel data interface and relies on an external clock and shutter trigger, making the module ideal for interfacing with an FPGA-based stereo imaging setup.

4.4 Camera Module Control Register

Bit	Bit Name	Bit Description	Default in Hex (Dec)	Shadowed	Legal Values (Dec)	Read/ Write
0x07 (7) Chip Control						
2:0	Scan Mode	0 = Progressive scan. 1 = Not valid. 2 = Two-field Interlaced scan. Even-numbered rows are read first, and followed by odd-numbered rows. 3 = Single-field Interlaced scan. If start address is even number, only even-numbered rows are read out; if start address is odd number, only odd-numbered rows are read out. Effective image size is decreased by half.	0	Y	0, 2, 3	W
3	Sensor Master/ Slave Mode	0 = Slave mode. Initiating exposure and readout is allowed. 1 = Master mode. Sensor generates its own exposure and readout timing according to simultaneous/ sequential mode control bit.	1	Y	0,1	W
4	Sensor Snapshot Mode	0 = Snapshot disabled. 1 = Snapshot mode enabled. The start of frame is triggered by providing a pulse at EXPOSURE pin. Sensor master/slave mode should be set to logic 1 to turn on this mode.	0	Y	0,1	W
5	Stereoscopy Mode	0 = Stereoscopy disabled. Sensor is stand-alone and the PLL generates a 320 MHz (x12) clock. 1 = Stereoscopy enabled. The PLL generates a 480 MHz (x18) clock.	0	Y	0,1	W
6	Stereoscopic Master/Slave mode	0 = Stereoscopic master. 1 = Stereoscopic slave. Stereoscopy mode should be enabled when using this bit.	0	Y	0,1	W
7	Parallel Output Enable	0 = Disable parallel output. DOUT(9:0) are in High-Z. 1= Enable parallel output.	1	Y	0,1	W
8	Simultaneous/ Sequential Mode	0 = Sequential mode. Pixel and column readout takes place only after exposure is complete. 1 = Simultaneous mode. Pixel and column readout takes place in conjunction with exposure.	1	Y	0,1	W

Table obtained from MT9V032 Datasheet [10]

4.5 Stereo Camera Schematic



4.6 Matlab Code

4.6.1 Camera Image Parsing

```
1 % Camera data parser - reads .log files from PuTTY for MT9V034 test
2 % Created by Georges Gauthier - 20 July 2016
3 clear all;
4 close all;
5
6 % prompt for a logfile; open selected file for reading
7 FILENAME = uigetfile('*.log','multiselect','off');
8 fprintf('File %s selected\n\r',FILENAME);
9 fid = fopen(FILENAME,'r');
10
11 image = zeros(480,752); % empty matrix that will hold final image
12 XPOS = 1; % current pixel X position
13 YPOS = 1; % current pixel Y position
14 LINENUM = 1; % number of pixels iterated through
15 ERRNUM = 0; % number of invalid pixels (happens when Tx is set too fast)
16
17 h = waitbar(0,'Parsing image...'); % show a loading bar
18 c = fgetl(fid); % get rid of 1st line
19
20 while 1 % iterate through the log file
21     c = fgetl(fid); % get the next line of the file
22     if ~ischar(c), break, end
23     if length(c) > 0 % if the given line contains valid data...
24         image(YPOS,XPOS) = str2num(c)/255; % ... store it as a pixel val
25     else % otherwise, throw an error
26         ERRNUM = ERRNUM + 1;
27         fprintf('Error #%d: Line %d contains no data\n\r',ERRNUM,LINENUM)
28     end
29     if XPOS<752 % update pixel x position
30         XPOS = XPOS + 1;
31     else % update pixel y position
32         XPOS = 1;
33         YPOS = YPOS + 1;
34     end
35     if mod(LINENUM,36096)==0 % update the loading bar every so often
36         waitbar(LINENUM /360000);
37     end
38     LINENUM = LINENUM + 1; % current line in file (for debug)
39 end
40
41 % display the image...
42 figure, imshow(image);
43 % ... also save the image to a file, overwrite if already saved
44 [path,name,ext] = fileparts(FILENAME);
45 imgname = strcat(name,'.png');
46 if (exist(imgname, 'file') == 2)
47     fprintf('File for image already exists... overwriting it\n\r')
48     delete(imgname);
```

```

49 end
50 saveas(gcf, imgname);
51 fprintf('Saved figure to image %s\n\r', imgname);
52
53 % close the file and waitbar before exit
54 fclose(fid);
55 close(h)

```

4.6.2 Disparity Algorithm Implementation

```

1 % The following code was adapted from a Mathworks example available here:
2 % http://www.mathworks.com/help/vision/examples/stereo-vision.html
3 %
4 % Original Revision by Chris McCormick
5 % http://mccormickml.com/2014/01/10/stereo-vision-tutorial-part-i/
6 %
7 % Modified by Georges Gauthier - glgauthier@wpi.edu
8 %
9 % This script will compute the disparity map for the image 'right.png' by
10 % correlating it to 'left.png' using basic block matching
11
12 clear all;
13 close all;
14
15 % Set to 1 to use 'Cones' Dataset
16 % Set to 0 to use your own image data (lines 26-29)
17 EXAMPLE_DATA = 1;
18
19 % Load the stereo images.
20 if (EXAMPLE_DATA == 0)
21     load('I1Rect.mat');
22     leftI = I1Rect;
23     load('I2Rect.mat');
24     rightI = I2Rect;
25 else
26     left = imread('left.png');
27     right = imread('right.png');
28     leftI = mean(left, 3);
29     rightI = mean(right, 3);
30 end
31
32 % DbasicSubpixel will hold the result of the block matching.
33 DbasicSubpixel = zeros(size(leftI), 'single');
34
35 % The disparity range defines how many pixels away from the block's
36 % location
37 % in the first image to search for a matching block in the other image.
38 % 50 appears to be a good value for the 450x375 images from the "Cones"
39 % dataset.
40 disparityRange = 50;

```

```

40
41 % Define the size of the blocks for block matching.
42 halfBlockSize = 5;
43 blockSize = 2 * halfBlockSize + 1;
44
45 % Get the image dimensions.
46 [imgHeight, imgWidth] = size(leftI);
47
48 % Create a progress bar
49 h = waitbar(0, 'Loading...');

50
51 % For each column 'm' of pixels in the image...
52 for (m = 1 : imgHeight)

53
54     % Set min/max row bounds for the template and blocks.
55     % e.g., for the first row, minr = 1 and maxr = 4
56     minr = max(1, m - halfBlockSize);
57     maxr = min(imgHeight, m + halfBlockSize);

58
59     % For each row 'n' of pixels in the image...
60     for (n = 1 : imgWidth)

61
62         % Set the min/max column bounds for the template.
63         % e.g., for the first column, minc = 1 and maxc = 4
64         minc = max(1, n - halfBlockSize);
65         maxc = min(imgWidth, n + halfBlockSize);

66
67         % Define the search boundaries as offsets from the template
68         % location.
69         % Limit the search so that we don't go outside of the image.
70         % 'mind' is the maximum number of pixels we can search to the
71         % left.
72         % 'maxd' is the maximum number of pixels we can search to the
73         % right.
74         %
75         % In the "Cones" dataset, we only need to search to the right, so
76         % mind
77         % is 0.
78         %
79         % For other images which require searching in both directions, set
80         % mind
81         % as follows:
82         % mind = max(-disparityRange, 1 - minc);
83         mind = 0;
84         maxd = min(disparityRange, imgWidth - maxc);

85
86         % Select the block from the right image to use as the template.
87         template = rightI(minr:maxr, minc:maxc);

88
89         % Get the number of blocks in this search.
90         numBlocks = maxd - mind + 1;

91
92         % Create a vector to hold the block differences.
93         blockDiffs = zeros(numBlocks, 1);

```

```

89
90    % Calculate the difference between the template and each of the
91    % blocks.
92    for (i = mind : maxd)
93
94        % Select the block from the left image at the distance 'i'.
95        block = leftI(minr:maxr, (minc + i):(maxc + i));
96
97        % Compute the 1-based index of this block into the 'blockDiffs'
98        % vector.
99        blockIndex = i - mind + 1;
100
101        % Take the sum of absolute differences (SAD) between the
102        % template
103        % and the block and store the resulting value.
104        blockDiffs(blockIndex, 1) = sum(sum(abs(template - block)));
105    end
106
107    % Sort the SAD values to find the closest match (smallest
108    % difference).
109    % Discard the sorted vector (the "~~" notation), we just want the
110    % list
111    % of indices.
112    [temp, sortedIndeces] = sort(blockDiffs);
113
114    % Get the 1-based index of the closest-matching block.
115    bestMatchIndex = sortedIndeces(1, 1);
116
117    % Convert the 1-based index of this block back into an offset.
118    % This is the final disparity value produced by basic block
119    % matching.
120    d = bestMatchIndex + mind - 1;
121
122    % Store the calculated disparity value in the resultant img matrix
123    DbasicSubpixel(m, n) = d;
124
125
126 end
127
128 % close the progress bar
129 close(h);
130
131 % Display the disparity map.
132 % Passing an empty matrix as the second argument tells imshow to take the
133 % minimum and maximum values of the data and map the data range to the
134 % display colors.
135 figure, imshow(DbasicSubpixel, []);

```

```

136 axis image;
137 colorbar;
138
139 % Specify the minimum and maximum values in the disparity map so that the
140 % values can be properly mapped into the full range of colors.
141 % If you have negative disparity values, this will clip them to 0.
142 caxis([0 disparityRange]);
143
144 % Set the title to display.
145 title(strcat('SAD Block Matching: ',num2str(blockSize),'x',...
146 num2str(blockSize), ' Block, ',num2str(disparityRange), 'px Search Range
147 '));
148
149 % plot both images in a final output graph
150 figure,
151 if (EXAMPLE_DATA == 0)
152 subplot(1,3,1), imshow(leftI)
153 title('Left Input Image')
154 subplot(1,3,3), imshow(rightI)
155 else
156 subplot(1,3,1), imshow(left)
157 title('Left Input Image')
158 subplot(1,3,3), imshow(right)
159 end
160 title('Right Input Image')
161 subplot(1,3,2), imshow(DbasicSubpixel,[])
162 title(strcat('SAD Block Matching: ',num2str(blockSize),'x',...
163 num2str(blockSize), ' Block, ',num2str(disparityRange), 'px Search Range
164 '));

```

4.7 Verilog Code

4.7.1 MT9V034 and Al422b Test Code

Top Module:

```
1  `timescale 1ns / 1ps
2  ///////////////////////////////////////////////////////////////////
3  // Created by Georges Gauthier
4  // July 09 2016
5  // Test module for controlling the Leopardboard LI-VM34LP camera breakout
6  ///////////////////////////////////////////////////////////////////
7  module mt9v034_top(
8      input sysclk, // 100MHz fpga clk
9      input reset, // addr2 reset
10     input cam_rst, // button for camera RESET_BAR
11     input trigger, // button for camera trigger
12     input SW_cam_oe, // switch for camera output enable
13     input cam_LV, // line valid in from camera
14     output LOCKED, // addr2 LOCKED led
15     output cam_sysclk, // sysclk out to camera
16     output cam_reset, // reset_bar out to camera
17     output cam_trigger, // trigger out to camera
18     output cam_oe, // output enable out to camera
19     output i2c_ready, // LED indicator for i2c bus ready
20     output [6:0] cathodes, // 7seg cathodes
21     output [3:0] anodes, // 7seg anodes
22     input MICRO_SW, // SW2, used to trigger a new FIFO dump
23         over UART from the mcs
24     input mcs_reset, // Microblaze reset
25     output MICRO_LEDO, // LED used to indicate if mcs is
26         reading from FIFO
27     input [7:0] FIFO_DATA, // D0[7:0] from AL422b fifo
28     output FIFO_WE, // Write enable to fifo (LV inverted)
29     output FIFO_OE, // read enable to fifo (active low)
30     output FIFO_RRST, // read reset to fifo (active low)
31     output FIFO_RCK, // rck to fifo (1MHz)
32     output UART_Tx // UART send pin from mcs
33 );
34
35     wire clk_20Hz_unbuf, clk_20Hz;
36     wire clk_10kHz;
37     wire clk_1MHz, clk_1MHz_unbuf;
38     wire clk_24MHz;
39     wire clk_100MHz;
40
41     // 24MHz clock for driving MT9V034's SYSCLK
42     // 100mhz out for FIFO
43     // note you can't connect sysclk to a dcm and other things
44     dcm CLK_24MHz
45     (
46         .CLK_IN1(sysclk),
47         .CLK_OUT1(clk_100MHz),
```

```

46     .CLK_OUT2(clk_24MHz),
47     .RESET(reset),
48     .LOCKED(LOCKED)
49   );
50
51 // further divide the dcm clock to other freqs
52 clk_div clks(
53   .reset(reset), // synchronous reset
54   .clk_24M(clk_24MHz), // 24MHz camera SCLK
55   .clk_fifo(clk_1MHz_unbuf), // 1MHz FIFO RCK
56   .clk_debounce(clk_20Hz_unbuf), // 20Hz clock pulse for debouncing
57   .stuff
58   .anodes(clk_10kHz) // 10k 7Seg anode driver
59 );
60
61 // clock buffer for 1MHz fifo rck
62 BUFG clk_1M (
63   .O(clk_1MHz),
64   .I(clk_1MHz_unbuf)
65 );
66 // clock buffer for 20Hz button debouncing
67 BUFG clk_20H (
68   .O(clk_20Hz),
69   .I(clk_20Hz_unbuf)
70 );
71 // forward the camera sysclk out using a dedicated clocking route
72 ODDR2 #(
73   .DDR_ALIGNMENT("NONE"), // Sets output alignment to "NONE", "C0" or
74   "C1"
75   .INIT(1'b0), // Sets initial state of the Q output to 1'b0 or 1'
76   .b1
77   .SRTYPE("SYNC") // Specifies "SYNC" or "ASYNC" set/reset
78 ) clkfwd0 (
79   .Q(cam_sysclk), // 1-bit DDR output data
80   .CO(clk_24MHz), // 1-bit clock input
81   .C1(~clk_24MHz), // 1-bit clock input
82   .CE(1'b1), // 1-bit clock enable input
83   .DO(1'b0), // 1-bit data input (associated with C0)
84   .D1(1'b1), // 1-bit data input (associated with C1)
85   .R(1'b0), // 1-bit reset input
86   .S(1'b0) // 1-bit set input
87 );
88
89 // forward the fifo read clk out using a dedicated clocking route
90 ODDR2 #(
91   .DDR_ALIGNMENT("NONE"), // Sets output alignment to "NONE", "C0" or
92   "C1"
93   .INIT(1'b0), // Sets initial state of the Q output to 1'b0 or 1'
94   .b1
95   .SRTYPE("SYNC") // Specifies "SYNC" or "ASYNC" set/reset
96 ) clkfwd1 (
97   .Q(FIFO_RCK), // 1-bit DDR output data
98   .CO(clk_1MHz), // 1-bit clock input

```

```

95      .C1(~clk_1MHz), // 1-bit clock input
96      .CE(1'b1), // 1-bit clock enable input
97      .DO(1'b0), // 1-bit data input (associated with C0)
98      .D1(1'b1), // 1-bit data input (associated with C1)
99      .R(1'b0), // 1-bit reset input
100     .S(1'b0) // 1-bit set input
101   );
102
103 // 7seg display controls
104 wire [15:0] displayVal;
105 seven_seg segs(
106   .values(displayVal), // values to be written to the four seven segment
107   // LEDs
108   .CLK(clk_24MHz), // 24MHz clock
109   .en(clk_10kHz), // 10kHz counter enable used for setting the
110   // segment refresh rate
111   .cathodes(cathodes),
112   .anodes(anodes)
113 );
114
115 // debounce trigger button input
116 debounce deb(
117   .clk(clk_20Hz),
118   .btn(trigger),
119   .btn_val(cam_trigger)
120 );
121
122 // debounce output enable switch
123 btnlatch sw_oe(
124   .clk(clk_20Hz),
125   .btn(SW_cam_oe),
126   .btn_val(cam_oe)
127 );
128
129 // debounce the microblaze input sw
130 wire read_en;
131 btnlatch fifoRead_en(
132   .clk(clk_20Hz),
133   .btn(MICRO_SW),
134   .btn_val(read_en)
135 );
136
137 // camera initialization sequence
138 reg [11:0] init_count = 12'h000;
139 always @(posedge clk_24MHz) // cam sysclk before ODDR2
140 begin
141   if (cam_RST) // if cam_RST is pressed, redo the initialization
142     sequence
143       init_count <= 12'h000;
144     else if(init_count < 2500) // keep cam_RST asserted for at least
145       20 cam_sysclk cycles - I use 30 since it's the minimum time for
146       the i2c bus to be ready
147       init_count <= init_count + 1'b1;
148 end

```

```

144 assign cam_reset = (init_count >= 20);
145 assign i2c_ready = (init_count >= 30);
146
147 // assert/de-assert RE and WE ~0.1mS after power on
148 wire fifo_rden;
149 assign FIFO_OE = fifo_rden;
150 assign FIFO_WE = ~cam_LV;
151
152 // Microblaze MCS for reading from local buffer/Tx over UART
153 wire fifo_read_en, fifo_reset; // tell fpga to put new data in the FIFO
154 wire [7:0] pixelVal; // value of a camera pixel from fpga line buffer ->
    microblaze
155 wire [9:0] pixelPos; // pixel position (0-751) on a line, from microblaze
    -> fpga line buffer
156 microblaze_mcs mcs_0 (
157     .Clk(clk_100MHz), // input Clk
158     .Reset(mcs_reset), // input Reset
159     .UART_Tx(UART_Tx), // output UART_Tx
160     .GP01({fifo_read_en,
161             fifo_reset,
162             MICRO0_LED0}), // output [3 : 0] GP01
163     .GP02(pixelPos),
164     .GPI1(pixelVal), // pixel data from FIFO/FPGA buffer
165     .GPI2({read_en,mcs_read_en}) // sw1
166 );
167
168 // Buffer for storing a line of pixels from the FIFO
169 fifo_read linebuf(
170     .reset_pointer(fifo_reset),
171     .get_data(fifo_read_en), // from microblaze (sent to trigger new
        read from FIFO to FPGA buffer)
172     .pixel_addr(pixelPos), // from microblaze, 0-751
173     .fifo_data(FIFO_DATA), // 8 bit data in from fifo
174     .fifo_rck(clk_1MHz), // 1MHz clock signal generated by FPGA
175     .fifo_rrst(FIFO_RRST), // fifo read reset (reset read addr pointer
        to 0)
176     .fifo_oe(fifo_rden), // fifo output enable (allow for addr pointer
        to increment)
177     .buffer_ready(mcs_read_en),
178     .pixel_value(pixelVal), // 8-bit pixel value from internal buffer
179     .current_line(displayVal)
180 );
181
182 endmodule

```

Local Data Buffer:

```

1 `timescale 1ns / 1ps
2 /////////////////////////////////
3 // Module for reading from the AL422b FIFO and storing pixel line data in
4 // local buffer.
5 /////////////////////////////////

```

```

6 module fifo_read(
7     input reset_pointer, // from microblaze, signal to assert
8         fifo_rrst
9     input get_data, // from microblaze (sent to trigger new read from
10        FIFO to FPGA buffer)
11     input [9:0] pixel_addr, // from microblaze, 0-751
12     input [7:0] fifo_data, // 8 bit data in from fifo
13     input fifo_rck, // 1MHz clock signal generated by FPGA
14     output reg fifo_rrst, // fifo read reset (reset read addr pointer
15        to 0)
16     output reg fifo_oe, // fifo output enable (allow for addr pointer
17        to increment)
18     output reg buffer_ready, // to microblaze, signal that buffer is
19        ready to read from
20     output [7:0] pixel_value, // 8-bit value from internal buffer
21     output [15:0] current_line // value to seven segment displays
22 );
23
24 parameter [1:0] ready = 2'b00;
25 parameter [1:0] read = 2'b01;
26 parameter [1:0] done = 2'b10;
27 parameter [1:0] init = 2'b11;
28
29 reg [1:0] state = ready;
30 reg [1:0] prev_state, next_state = ready;
31
32 reg [7:0] pixel_line [0:751]; // implemented in BRAM
33 reg [9:0] pixel = 10'b000_0000_0000;
34 reg [15:0] num_lines = 16'h0000;
35
36 always @(posedge fifo_rck)
37     state <= next_state;
38
39 always @(state,get_data,pixel)
40     case(state)
41         ready:
42             begin
43                 if(get_data)
44                     next_state = init;
45                 else
46                     next_state = ready;
47
48                 prev_state = ready;
49             end
50         init:
51             begin
52                 next_state = read;
53                 prev_state = init;
54             end
55         read:
56             begin
57                 if(pixel == 751)
58                     next_state = done;
59                 else
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
589

```

```

55                               next_state = read;
56
57                               prev_state = read;
58                           end
59           done:
60               begin
61                   next_state = ready;
62                   prev_state = done;
63               end
64           endcase
65
66 always @ (posedge fifo_rck)
67 begin
68     if (reset_pointer)
69         begin
70             fifo_rrst <= 1'b0;
71             num_lines <= 16'h0000;
72         end
73     else if (state == ready) // allow for MCS to read from
74         pixel_line
75         begin
76             //pixel_value [7:0] <= pixel_line[pixel_addr];
77             fifo_rrst <= 1'b1; // make sure read addr doesn't
78                 get reset
79         end
80     else if (state == init) // prepare to read new data from
81         the AL422 into pixel_line
82         begin
83             pixel <= 10'b00_0000_000;
84             num_lines <= num_lines + 1'b1;
85             buffer_ready <= 1'b0;
86             fifo_oe <= 1'b0; // allow for read pointer to
87                 increment
88         end
89     else if (state == read) // read data in from the AL422
90         begin
91             if (next_state == done)
92                 fifo_oe <= 1'b1; // turn off read enable
93             if (prev_state != init) // one cycle delay between
94                 init and valid data
95             begin
96                 pixel_line[pixel] <= fifo_data;
97                 pixel <= pixel + 1'b1;
98             end
99         end
100    else if (state == done)
101        begin
102            buffer_ready <= 1'b1;
103        end
104
105 // display number of lines written on 7seg display
106 assign current_line = (num_lines);
107 // allow for MCS to read stored pixel line at given addr if state==ready

```

```
104 assign pixel_value [7:0] = pixel_line[pixel_addr];  
105  
106 endmodule
```

4.8 C Code

4.8.1 Camera Image Parsing

```
1  /*
2   * Source code for printing values from the AL422B FIFO / FPGA Buffer over
3   *      UART
4   */
5
6 #include <stdio.h>
7 #include "platform.h"
8 #include "xparameters.h"
9 #include "xiomodule.h"
10
11 // GPO1
12 #define GETDATA (1<<2) // load a new line of pixels into the FPGA buffer
13 #define RRST (1<<1) // reset to address 0
14 #define LED (1<<0) // LED indicator
15 // GPIO2
16 #define SW_READ (1<<1)
17 #define BUF_READY (1<<0)
18
19 void print(char *str);
20 void _EXFUN(xil_printf, (const char*, ...));
21
22 int main()
23 {
24     init_platform();
25     int pixel_position = 0, row = 0;;
26     u8 data=0x00, GP01=0x00, GPI2=0x00, swState=0x00, prevState=0x00;
27
28     XIOModule gpi;
29     XIOModule gpo;
30
31     // GPIO1 = pixel_value(7:0)
32     XIOModule_Initialize(&gpi, XPAR_IOMODULE_0_DEVICE_ID);
33     XIOModule_Start(&gpi);
34
35     // GP01 = (GETDATA)|(RRST)|(LED)
36     XIOModule_Initialize(&gpo, XPAR_IOMODULE_0_DEVICE_ID);
37     XIOModule_Start(&gpo);
38
39     print("\n\rMT9V034 controller and AL422B FIFO reader\n\r");
40     while(1)
41     {
42         // get switch position
43         GPI2 = XIOModule_DiscreteRead(&gpi, 2);
44
45         if((GPI2&SW_READ)!=0){
46             swState = 1;
47             if (row >= 480) GP01 &= ~(LED);
```

```

48     else GPO1 |= LED;
49     GPO1 &= ~(RRST|GETDATA);
50     XIOModule_DiscreteWrite(&gpo,1, GPO1);
51 }else{
52     GPO1 &= ~(RRST|LED|GETDATA);
53     XIOModule_DiscreteWrite(&gpo,1, GPO1);
54     row = 0;
55     swState = 0;
56 }
57
58 // code below runs only once based on SW state change
59 if (prevState != swState){
60     if(swState){
61         print("\n\rReading from FIFO...\n\r");
62
63         GPO1 |= (RRST); // reset FIFO position to 0th index
64         GPO1 &=~ (GETDATA); // make sure we're not trying to read
65             data
66         XIOModule_DiscreteWrite(&gpo,1, GPO1);
67         pixel_position = 0;
68         GPO1 &=~(RRST);
69         XIOModule_DiscreteWrite(&gpo,1, GPO1);
70         GPO1 |= (GETDATA); // make sure we're not trying to read
71             data
72         XIOModule_DiscreteWrite(&gpo,1, GPO1);
73
74         u32 pixelsRead = 0;
75
76         while(row<480){
77             // make sure read_sw hasn't been turned off
78             GPI2 = XIOModule_DiscreteRead(&gpi,2);
79             if ((GPI2&SW_READ)==0) break;
80
81             u8 i=0;
82             // check to see if BUF_READY is good to go
83             GPI2 = XIOModule_DiscreteRead(&gpi,2);
84             // wait until it is
85             while((GPI2&BUF_READY)==0){
86                 if(i==0){
87                     i++;
88                     //print("\n\t buffer not ready \n\r");
89                 }
90                 GPI2 = XIOModule_DiscreteRead(&gpi,2);
91             }
92             GPO1 &=~ (GETDATA); // make sure we're not trying to
93                 read data
94             XIOModule_DiscreteWrite(&gpo,1, GPO1);
95
96             while (pixel_position < 752){
97                 // update pixel position for FPGA buffer
98                 XIOModule_DiscreteWrite(&gpo,2,pixel_position);
99
100                // make sure read_sw hasn't been turned off
101                GPI2 = XIOModule_DiscreteRead(&gpi,2);

```

```

99         if ((GPIO2&SW_READ)==0) break;
100
101        // read value at pixel position from FPGA buffer
102        data = XIOModule_DiscreteRead(&gpi,1);
103
104        //print the value
105        xil_printf("%d\n\r",data);
106
107        // increment to the next pixel position
108        pixel_position++;
109        pixelsRead++;
110    }
111    // signal to the FPGA that we want more data!
112    GP01 |= (GETDATA);
113    XIOModule_DiscreteWrite(&gpo,1,GP01);
114    pixel_position = 0;
115    row++;
116    //xil_printf("Row: %d",row);
117
118    }
119    //xil_printf("%d Pixels Read by MCS",pixelsRead);
120 } else {
121    GP01 &= ~(GETDATA);
122    GP01 |= RRST;
123    XIOModule_DiscreteWrite(&gpo,1,GP01);
124    print("\n\rReset for new sequence\n\r");
125    }
126
127    prevState = swState; // update prev switch position
128 }
129 cleanup_platform();
130 return 0;
132 }
```

4.9 LaTeX Coding Examples

This section isn't intended to remain here, but can serve as an example for how to set things up later on

4.9.1 Figures



Figure 29: A Test Figure

Using the \ref command, I'm able to reference Figure 29 by calling \ref{wpiLogo}.

4.9.2 Code Snippet

Code snippets can be created by calling \begin{lstlisting}[style=<Language_Name>], inserting all code, and then calling \end{lstlisting}. Also call \singespacing before the code snippet and \doublespacing after to keep things from getting too big.

Note that you can also use

```
\lstinputlisting[src.ext][style=<Language_Name>,firstline=lineNumber, lastline=lineNumber]
```

with reference to source files and listings will import the code for you. Way less work!

```
1 //verilog code example
2 always @ (x, y, z)
3   x <= y + z;
```

```
1 %Matlab code example
2 clear all;
3 close all;
4 FILENAME = uigetfile('*.log','multiselect','off');
5 fprintf('File %s selected\n\r',FILENAME);
6 return;
```

```
1 //C code example
2 #include "someheader.h"
```

```
3 XIOModule gpo;
4
5 void GPOWrite(u8 value){
6     XIOModule_DiscreteWrite(&gpo, 1, value);
7     xil_printf("Wrote %d to GPO\n\r", value);
8     return;
9 }
```

4.9.3 Using the bibliography

All bibliographic references are contained in `bib.tex`. To cite a reference in the paper, use the `\cite` command.

As an example, I can cite *Serveball* at the end of this sentence by calling `\cite{serveball}.`[6]

To cite multiple references, call `\cite{ref1,ref2}.`[6, 12]